

1 **POINT SPREAD FUNCTION APPROXIMATION OF HIGH RANK**
2 **HESSIANS WITH LOCALLY SUPPORTED NON-NEGATIVE**
3 **INTEGRAL KERNELS***

4 NICK ALGER[†], TUCKER HARTLAND[‡], NOEMI PETRA[§], AND OMAR GHATTAS[¶]

5 **Abstract.** We present an efficient matrix-free point spread function (PSF) method for approximating operators that have locally supported non-negative integral kernels, with particular focus
6 on approximating high rank Hessians in distributed parameter inverse problems governed by partial
7 differential equations (PDEs). The proposed method uses impulse responses of the underlying operator
8 (computed at a collection of scattered points) and interpolates translated and scaled versions
9 of these impulse responses to approximate entries of the integral kernel. The interpolation uses a
10 “local mean displacement invariance” approximation, which improves upon classical local translation
11 invariance. Impulse responses are computed by applying the operator to a small number of
12 Dirac combs associated with “batches” of point sources. By solving an ellipsoid packing problem,
13 we choose as many point sources as possible per batch, while ensuring that the supports of the
14 impulse responses within each batch do not overlap. Support ellipsoids are estimated a-priori based
15 on impulse response moments, and these moments are computed for impulse responses at all points
16 simultaneously via a procedure that involves applying the operator to a small number of polynomial
17 functions. The ability to rapidly evaluate kernel entries allows us to construct a hierarchical matrix
18 (H-matrix) approximation of the operator. Fast H-matrix methods are used to perform further matrix
19 computations in nearly linear complexity, including matrix-matrix addition and multiplication,
20 matrix inversion and factorization, and matrix-vector products. We illustrate our approach by applying
21 the method to approximate the Hessians in an ice sheet flow PDE constrained inverse problem
22 with highly informative data, and an advection dominated advection-diffusion PDE constrained inverse
23 problem. Numerical results demonstrate that the proposed method substantially outperforms
24 existing methods. We show that our method can approximate high rank operators using only a small
25 number of operator applies.
26

27 **Key words.** data scalability, Hessian, hierarchical matrix, matrix-free, operator approximation,
28 PDE constrained inverse problems, point spread function, product convolution, local translation
29 invariance, impulse response, moment methods

30 **AMS subject classifications.** 35R30, 41A05, 41A35, 41A36, 42A85, 47A52, 47A58, 49K20,
31 65D12, 65F08, 65F10, 65N20, 65N21, 86A22, 86A40

32 **1. Introduction.** We present an efficient *matrix-free* point spread function (PSF)
33 method for approximating operators, $\mathcal{A} : L^2(\Omega) \rightarrow L^2(\Omega)'$, that have locally sup-
34 ported non-negative integral kernels. Here $\Omega \subset \mathbb{R}^d$ is a bounded domain, and $L^2(\Omega)'$
35 is the space of real-valued continuous linear functionals on $L^2(\Omega)$. We focus on ap-
36 proximating Hessians in optimization and inverse problems governed by partial dif-
37 ferential equations (PDEs) [5, 44], though such operators also arise as certain Schur
38 complements in Schur complement methods for solving partial differential equations
39 (PDEs) [50, 64], Poincare-Steklov operators in domain decomposition methods (e.g.,

*Submitted to the editors DATE.

Funding: This research was supported by the National Science Foundation under Grant No. DMS-1840265 and CAREER-1654311, DOD grant FA9550-21-1-0084, and DOE grants DE-SC0021239 and DE-SC0019303.

[†]Oden Institute for Computational Engineering and Sciences, The University of Texas at Austin, Austin, TX (nalger@oden.utexas.edu).

[‡]Department of Applied Mathematics, University of California, Merced, Merced, CA. (thardtland@ucmerced.edu).

[§]Department of Applied Mathematics, University of California, Merced, Merced, CA. (npe-tra@ucmerced.edu).

[¶]Oden Institute for Computational Engineering and Sciences and Walker Department of Mechanical Engineering, The University of Texas at Austin, Austin, TX (omar@oden.utexas.edu).

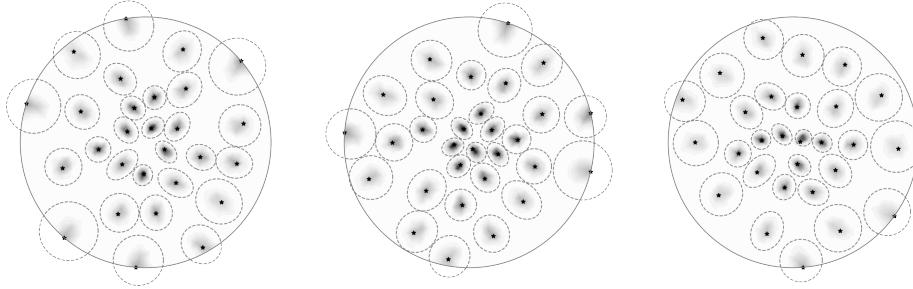


Fig. 1: Batches, η_b , of normalized impulses, ϕ_x , for the Stokes inverse problem data misfit Gauss-Newton Hessian (Section 7). Black stars are point source locations. Dashed gray ellipses are estimated impulse response support ellipsoids. The large circle is $\partial\Omega$.

40 Dirichlet-to-Neumann maps), covariance operators in spatial statistics [14, 31, 53],
 41 and blurring operators in imaging [17, 58]. Here “matrix-free” means that we may
 42 apply \mathcal{A} and its transpose¹, \mathcal{A}^T , to arbitrary functions,

43 (1.1)
$$u \mapsto \mathcal{A}u \quad \text{and} \quad w \mapsto \mathcal{A}^T w,$$

44 via a black box computational procedure, but we cannot easily access entries of \mathcal{A} 's
 45 integral kernel. Evaluating the maps in (1.1) may require, as a sub procedure, per-
 46 forming a costly computation, such as solving a large linear system or time stepping.

47 We use *impulse response interpolation* to form a high rank operator approximation
 48 using a small number of operator applies. The impulse response, ϕ_x , associated with a
 49 point x is the Riesz representation² of the linear functional that results from applying
 50 \mathcal{A} to a delta distribution (i.e., point source, impulse) centered at x . We compute
 51 batches of impulse responses by applying \mathcal{A} to weighted sums of delta distributions
 52 associated with batches of points scattered throughout the domain (see Figure 1).
 53 Then we interpolate translated and scaled versions of these impulse responses to
 54 approximate entries of the operator's integral kernel. Picking the batches of points
 55 requires us to estimate the supports of the impulse responses ϕ_x before we compute
 56 them. The idea of estimating the supports of the functions ϕ_x a-priori was inspired
 57 by techniques from resolution analysis in seismic imaging. There, \mathcal{A}^T is applied to
 58 a random noise function, and the width of ϕ_x is estimated to be the autocorrelation
 59 length of the resultant function near x [26, 68]. We use polynomial functions instead
 60 of random noise functions (see Section 4.1). Our method estimates the support of ϕ_x
 61 more accurately and reliably than random noise probing at the cost of the additional
 62 constraint that \mathcal{A} has a non-negative integral kernel. Our method may be categorized
 63 as a PSF method that is loosely based on “product convolution” (PC) approximations,
 64 which are approximations of an operator by weighted sums of convolution operators
 65 with spatially varying weights. PC and PSF methods have a long history dating back
 66 several decades. We note the following papers (among many others), [1, 5, 10, 22],

¹Recall that $\mathcal{A}^T : L^2(\Omega) \rightarrow L^2(\Omega)'$ is the unique operator satisfying $(\mathcal{A}u)(w) = (\mathcal{A}^Tw)(u)$ for all $u, w \in L^2(\Omega)$, where $\mathcal{A}u \in L^2(\Omega)'$ is the result of applying \mathcal{A} to $u \in L^2(\Omega)$, and $(\mathcal{A}u)(w)$ is the result of applying that linear functional to $w \in L^2(\Omega)$, and similar for operations with \mathcal{A}^T .

²Recall that the Riesz representative of a functional $\rho \in L^2(\Omega)'$ with respect to the L^2 inner product is the unique function $\rho^* \in L^2(\Omega)$ such that $\rho(w) = (\rho^*, w)_{L^2(\Omega)}$ for all $w \in L^2(\Omega)$.

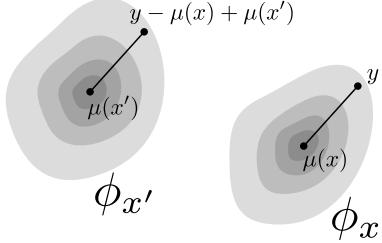


Fig. 2: \mathcal{A} is locally mean displacement invariant if $\phi_x(y) \approx \phi_{x'}(y - \mu(x) + \mu(x'))$ when x is close to x' . Here $\mu(z)$ denotes the mean (center of mass) of ϕ_z .

67 [24, 25, 27, 58, 72], in which the convolution kernels are constructed from sampling
 68 impulse responses of the operator to scattered point sources. For background on
 69 PC and PSF methods, we recommend the following papers: [18, 23, 30]. While PC
 70 approximations are based on an assumption of local translation invariance, the method
 71 we propose is based on an assumption we call “local mean displacement invariance,”
 72 which is more general than local translation invariance, and includes both CUR low
 73 rank approximation and local translation invariance as special cases (Section 4.2 and
 74 Figure 2).

75 The ability to rapidly approximate entries of \mathcal{A} ’s integral kernel allows us to form
 76 a hierarchical matrix (H-matrix) [11, 38] approximation of a discretized version of
 77 \mathcal{A} . H-matrices are a matrix format in which the rows and columns of the matrix are
 78 re-ordered, then the matrix is recursively subdivided into blocks, in such a way that
 79 many off-diagonal blocks are low rank, even though the matrix as a whole may be high
 80 rank. H-matrix methods permit us to perform matrix-vector products cheaply, and
 81 perform other useful linear algebra operations that cannot be done easily using the
 82 original operator. These operations include matrix-matrix addition, matrix-matrix
 83 multiplication, matrix factorization, and matrix inversion. The work and memory
 84 required to perform these operations for an $N \times N$ H-matrix scales *nearly linearly*
 85 in N (i.e., $o(N^{1+\epsilon})$ for any $\epsilon > 0$). The exact cost depends on the type of H-matrix
 86 used, the operation being performed, and the rank of the off-diagonal blocks [35]. The
 87 ability to perform these operations permits, for example, fast solution of Newton linear
 88 systems in PDE constrained optimization, fast sampling of ill-conditioned posterior
 89 distribution in Bayesian inverse problems, and construction of high rank Gaussian
 90 surrogate models that can be used for uncertainty quantification.

91 **2. Why we need high rank Hessian approximations.** In distributed pa-
 92 rameter inverse problems governed by PDEs, one seeks to infer an unknown spatially
 93 varying parameter field from limited observations of a state variable that depends on
 94 the parameter implicitly through the solution of a PDE. Conventionally, the inverse
 95 problem is formulated using either a deterministic framework [8, 70], or a Bayesian
 96 probabilistic framework [46, 66, 67]. In the deterministic framework, one solves an
 97 optimization problem to find the parameter that best fits the observations, subject to
 98 appropriate regularization [20]. In the probabilistic framework, Bayes’ theorem com-
 99 bines the observations with prior information to form a posterior distribution over
 100 the space of all possible parameter fields, and computations are performed to extract
 101 statistical information about the parameter from this posterior. The Hessian of the
 102 objective function in the deterministic optimization problem and the Hessian of the

103 negative log posterior in the Bayesian setting are equal or approximately equal to
 104 each other under typical noise, regularization, and prior models, so we refer to both
 105 of these Hessians as “the Hessian.” The Hessian consists of a data misfit term (the
 106 “data misfit Hessian”), which depends on discrepancy between the observations and
 107 the associated model predictions, and a regularization or prior term (the “regulariza-
 108 tion Hessian”) which does not depend on the observations. For more details on the
 109 Hessian, see [4, 32, 69].

110 Hessian approximations and preconditioners are highly desirable because the Hes-
 111 sian is central to efficient solution of inverse problems in both the deterministic and
 112 Bayesian settings. When solving the deterministic optimization problem with Newton-
 113 type methods, the Hessian is the coefficient operator for the linear system that must
 114 be solved or approximately solved at every Newton iteration. Good Hessian precondi-
 115 tioners reduce the number of iterations required to solve these Newton linear systems
 116 with Krylov methods such as conjugate gradient [63]. In the Bayesian setting, the
 117 inverse of the Hessian is the covariance of a local Gaussian approximation of the
 118 posterior. Such a Gaussian distribution can be used as a direct approximations of
 119 the posterior, or as a proposal for Markov chain Monte-Carlo methods for drawing
 120 samples from the posterior. For instance, see [47, 60] and the references therein.

121 Owing to the implicit dependence of predicted observations on the parameter,
 122 entries of the Hessian are not easily accessible. Rather, the Hessian may be applied to
 123 a vector via a computational process that involves solving a pair of forward and adjoint
 124 PDEs which are linearizations of the original PDE [32, 37]. The most popular matrix-
 125 free Hessian approximation methods are based on low rank approximation of either the
 126 data misfit Hessian, or the data misfit Hessian preconditioned by the regularization
 127 Hessian, e.g., [13, 16, 28, 60, 65]. Krylov methods such as Lanczos or randomized
 128 methods [15, 41] are typically used to construct these low rank approximations by
 129 applying the Hessian to vectors. Using these methods, the required number of Hessian
 130 applications (and hence the required number of PDE solves) is typically proportional
 131 to the rank of the low rank approximation. Low rank approximation methods are
 132 justified by arguing that the numerical rank of the data misfit Hessian is insensitive
 133 to the dimension of the discretized parameter. This means that the required number
 134 of PDE solves remains the same as the mesh used to discretize the parameter is
 135 refined. However, in many inverse problems of practical interest the numerical rank
 136 of the data misfit Hessian, while mesh independent, is still large, which makes it costly
 137 to approximate the Hessian using low rank approximation methods [7, 13, 44].

138 Examples of inverse problems with high rank data misfit Hessians include ad-
 139 vection dominated advection-diffusion inverse problems [2][29, Chapter 5], high fre-
 140 quency wave inverse problems [13], and more generally, all inverse problems in which
 141 the observations highly inform the parameter. The eigenvalues of the data misfit
 142 Hessian characterize the sensitivity of the predicted observations to perturbations of
 143 the parameter in the corresponding eigenvector directions, so more informative data
 144 leads to larger eigenvalues and therefore a larger numerical rank [3][4, Section 1.4
 145 and Chapter 4]. Roughly speaking, the numerical rank of the data misfit Hessian
 146 is the dimension of the subspace of parameter space that is informed by the data.
 147 The numerical rank of the regularization preconditioned data misfit Hessian may be
 148 reduced by increasing the strength of the regularization, but this throws away useful
 149 information—components of the parameter that could be learned from components of
 150 the observations that have a large signal to noise ratio are instead reconstructed based
 151 on the regularization [6, Section 4][70, Chapters 1 and 7]. Hence, low rank approxi-
 152 mation methods suffer from a predicament—if the data highly inform the parameter

and the regularization is chosen appropriately, then a large number of operator applies are required to form an accurate approximation of the Hessian using low rank approximation methods. High rank Hessian approximation methods are needed.

Recently there have been improvements in matrix-free H-matrix construction methods in which an operator is applied to structured random vectors, and the response of the operator to those random vectors is processed to construct an H-matrix approximation [51, 52, 55, 56, 57]. These methods (which we do not use here) have been used to approximate Hessians in PDE constrained inverse problems [7, 42]. Although the required number of matrix-vector products grows nearly linearly in the discretized parameter dimension and linearly in the rank of the off-diagonal blocks, the required number of matrix-vector products is still large (e.g., hundreds to thousands). In this paper, we also form an H-matrix Hessian approximation. However, to reduce the required number of operator actions, we first form a PSF approximation of the data misfit Hessian by exploiting locality and non-negative integral kernel properties, then form the H-matrix using classical H-matrix techniques. Not all data misfit Hessians satisfy the local non-negative integral kernel properties (for example, the wave inverse problem data misfit Hessian has a substantial amount of negative entries in its integral kernel). However, many data misfit Hessians of practical interest do satisfy these properties (either exactly or approximately), and our method is targeted at approximating these Hessians.

3. Preliminaries. Let $\Omega \subset \mathbb{R}^d$ be a bounded domain (typically $d = 1, 2$, or 3). We seek to approximate integral operators $\mathcal{A} : L^2(\Omega) \rightarrow L^2(\Omega)'$ of the form

$$(3.1) \quad (\mathcal{A}u)(w) := \int_{\Omega} \int_{\Omega} w(y) \Phi(y, x) u(x) dx dy.$$

Here the linear functional $\mathcal{A}u \in L^2(\Omega)'$ is the result of applying \mathcal{A} to $u \in L^2(\Omega)$, the scalar $(\mathcal{A}u)(w)$ is the result of applying that linear functional to $w \in L^2(\Omega)$. The integral kernel, $\Phi : \Omega \times \Omega \rightarrow \mathbb{R}$, exists in principle but is not easily accessible. In this section we describe how to extend the action of \mathcal{A} to distributions, which allows us to define impulse responses (Section 3.1), we state the conditions on \mathcal{A} that our method requires (Section 3.2), and we detail finite element discretization (Section 3.3).

3.1. Distributions and impulse responses. The action of \mathcal{A} may be extended to distributions if Φ is sufficiently regular. Let $\rho \in L^2(\Omega)'$, and let $\rho^* \in L^2(\Omega)$ denote the Riesz representative of ρ with respect to the $L^2(\Omega)$ inner product. We have

$$(3.2a) \quad (\mathcal{A}\rho^*)(w) = \int_{\Omega} \int_{\Omega} w(y) \Phi(y, x) \rho^*(x) dx dy$$

$$(3.2b) \quad = \int_{\Omega} w(y) \int_{\Omega} \Phi(y, x) \rho^*(x) dx dy = \int_{\Omega} w(y) \rho(\Phi(y, \cdot)) dy,$$

where $\Phi(y, \cdot)$ denotes the function $x \mapsto \Phi(y, x)$. Now let $\mathcal{D}(\Omega) \subset L^2(\Omega)$ be a suitable space of test functions and let $\rho : \mathcal{D}(\Omega) \rightarrow \mathbb{R}$ be a distribution. In this case, the Riesz representative ρ^* may not exist, so the derivation in (3.2) is not valid. However, if Φ is sufficiently regular such that the function $y \mapsto \rho(\Phi(y, \cdot))$ is well defined for almost all $y \in \Omega$, and if this function is in $L^2(\Omega)$, then the right hand side of (3.2b) is well defined. Hence, we *define* the action of \mathcal{A} on the distribution ρ to be the right hand side of (3.2b). We denote this action by “ $\mathcal{A}\rho^*$,” even if ρ^* does not exist.

Let δ_x denote the delta distribution³ (i.e., point source, impulse) centered at the

³Recall that the delta distribution $\delta_x : \mathcal{D}(\Omega) \rightarrow \mathbb{R}$ is defined by $\delta_x(w) = w(x)$ for all $w \in \mathcal{D}(\Omega)$.

196 point $x \in \Omega$. The *impulse response* of \mathcal{A} associated with x is the function $\phi_x : \Omega \rightarrow \mathbb{R}$,

197 (3.3)
$$\phi_x := (\mathcal{A}\delta_x^*)^*,$$

198 that is formed by applying \mathcal{A} to δ_x , then taking the Riesz representation of the
199 resulting linear functional. Using (3.2b) and the definition of the delta distribution,
200 we see that ϕ_x may also be written as the function $\phi_x(y) = \Phi(y, x)$.

201 **3.2. Required conditions.** We focus on approximating operators that satisfy
202 the following conditions:

- 203 1. The kernel Φ is sufficiently regular so that ϕ_x is well defined for all $x \in \Omega$.
204 2. The supports of the impulse responses ϕ_x are contained in localized regions.
205 3. The integral kernel is non-negative in the sense that

$$\Phi(y, x) \geq 0$$

205 for all $(y, x) \in \Omega \times \Omega$.⁴

206 Our method may still perform well if these conditions are relaxed slightly. It is fine if
207 the support of ϕ_x is not perfectly contained in a localized region, as long as the bulk of
208 the “mass” of ϕ_x is contained in a localized region. The integral kernel does not need
209 to be non-negative for all pairs of points $(y, x) \in \Omega \times \Omega$, as long as it is non-negative
210 for the vast majority of pairs of points (y, x) , and as long as the negative numbers are
211 comparatively small. If these conditions are violated, our method will incur additional
212 error. If these conditions are violated too much, our method may fail.

213 **3.3. Finite element discretization.** In computations, functions are discretized
214 and replaced by finite dimensional vectors, and operators mapping between infinite
215 dimensional spaces are replaced by operators mapping between finite dimensional spaces.
216 In this paper we discretize using continuous Galerkin finite elements satisfying the
217 Kronecker property (defined below). With minor modifications, our method could be
218 used with more general finite element methods, or other discretization schemes such
219 as finite differences or finite volumes.

220 Let $\psi_1, \psi_2, \dots, \psi_N$ be a set of continuous Galerkin finite element basis functions
221 used to discretize the problem on a mesh with mesh size parameter h , let $V_h :=$
222 $\text{span}(\psi_1, \psi_2, \dots, \psi_N)$ be the corresponding finite element space under the L^2 inner
223 product, and let $\zeta_i \in \mathbb{R}^d$, $i = 1, \dots, N$ be the Lagrange nodes associated with the
224 functions ψ_i . We assume that the finite element basis satisfies the Kronecker property,
225 i.e., $\psi_i(\zeta_i) = 1$ and $\psi_i(\zeta_j) = 0$ if $i \neq j$. For $u_h \in V_h$ we write $\mathbf{u} \in \mathbb{R}_M^m$ to denote
226 the coefficient vector for u_h with respect to the finite element basis, i.e., $u_h(x) =$
227 $\sum_{i=1}^N \mathbf{u}_i \psi_i(x)$. Linear functionals $\rho_h \in V'_h$ have coefficient dual vectors $\boldsymbol{\rho} \in \mathbb{R}_{M^{-1}}^m$,
228 with entries $\boldsymbol{\rho}_i = \rho_h(\psi_i)$ for $i = 1, \dots, m$. Here $M \in \mathbb{R}^{N \times N}$ denotes the sparse finite
229 element mass matrix which has entries $M_{ij} = \int_{\Omega} \psi_i(x) \psi_j(x) dx$ for $i, j = 1, \dots, N$.
230 The space \mathbb{R}_M^N is \mathbb{R}^N with the inner product $(\mathbf{u}, \mathbf{w})_M := \mathbf{u}^T M \mathbf{w}$, and $\mathbb{R}_{M^{-1}}^N$ is the
231 analogous space with M^{-1} replacing M . Direct calculation shows that \mathbb{R}_M^N and $\mathbb{R}_{M^{-1}}^N$
232 are isomorphic to V_h and V'_h as Hilbert spaces, respectively.

233 After discretization, the operator $\mathcal{A} : L^2(\Omega) \rightarrow L^2(\Omega)'$ is replaced by an operator
234 $A_h : V_h \rightarrow V'_h$, which becomes an operator

235
$$\mathbf{A} : \mathbb{R}_M^N \rightarrow \mathbb{R}_{M^{-1}}^N$$

⁴Note that having a non-negative integral kernel is different from positive semi-definiteness. The operator \mathcal{A} need not be positive semi-definite to use our method, and positive semi-definite operators need not have a non-negative integral kernel.

under the isomorphism discussed above. Our method is agnostic to the computational procedure for approximating \mathcal{A} with \mathbf{A} . What is important is that we do not have direct access to matrix entries \mathbf{A}_{ij} . Rather, we have a computational procedure that allows us to compute matrix-vector products $\mathbf{u} \mapsto \mathbf{Au}$ and $\mathbf{w} \mapsto \mathbf{A}^T\mathbf{w}$, and computing these matrix-vector products is costly. Of course, matrix entries can be computed via matrix-vector products as $\mathbf{A}_{ij} = (\mathbf{A}\mathbf{e}_j)_i$, where $\mathbf{e}_j = (0, \dots, 0, 1, 0, \dots, 0)^T$ is the length N unit vector with one in the j th coordinate and zeros elsewhere. But computing the matrix-vector product $\mathbf{e}_j \mapsto \mathbf{A}\mathbf{e}_j$ is costly, and therefore wasteful if we do not use other matrix entries in the j th column of \mathbf{A} . Hence, methods for approximating \mathbf{A} are computationally intractable if they require accessing scattered matrix entries from many different rows and columns of \mathbf{A} .

The operator $A_h : V_h \rightarrow V'_h$ can be written in integral kernel form, (3.1), but with Φ replaced by a slightly different integral kernel, Φ_h , which we do not know, and which differs from Φ due to discretization error. Since the functions in V_h are continuous at x , the delta distribution δ_x is a continuous linear functional on V_h , which has a discrete dual vector $\boldsymbol{\delta}_x \in \mathbb{R}_{\mathbf{M}^{-1}}^N$ with entries $(\boldsymbol{\delta}_x)_i = \psi_i(x)$ for $i = 1, \dots, N$. Additionally, it is straightforward to verify that the Riesz representation, $\rho_h^* \in V_h$, of a functional $\rho \in V'_h$ has coefficient vector $\boldsymbol{\rho}^* = \mathbf{M}^{-1}\boldsymbol{\rho}$. Therefore, the formula for the impulse response from (3.3) becomes $\phi_x = (A_h\delta_x^*)^* = \mathbf{M}^{-1}\mathbf{A}\mathbf{M}^{-1}\boldsymbol{\delta}_x$, and the (y, x) kernel entry of Φ_h may be written as $\Phi_h(y, x) = \boldsymbol{\delta}_y^T\phi_x = \boldsymbol{\delta}_y^T\mathbf{M}^{-1}\mathbf{A}\mathbf{M}^{-1}\boldsymbol{\delta}_x$. Now define $\boldsymbol{\Phi} \in \mathbb{R}^{N \times N}$ to be the following dense matrix of kernel entries evaluated at all pairs of Lagrange nodes:

$$(3.4) \quad \Phi_{ij} := \Phi_h(\zeta_i, \zeta_j).$$

Because of the Kronecker property of the finite element basis, we have $\boldsymbol{\delta}_{\zeta_i} = \mathbf{e}_i$. Thus, we have $\Phi_h(\zeta_i, \zeta_j) = (\mathbf{M}^{-1}\mathbf{A}\mathbf{M}^{-1})_{ij}$, which implies

$$(3.5) \quad \mathbf{A} = \mathbf{M}\boldsymbol{\Phi}\mathbf{M}.$$

Broadly, our method will construct an H-matrix approximation of \mathbf{A} by forming H-matrix approximations of $\boldsymbol{\Phi}$ and \mathbf{M} (or a lumped mass version of \mathbf{M}), then multiplying these matrices per (3.5) using H-matrix methods. Classical H-matrix construction methods require access to arbitrary matrix entries Φ_{ij} , but these matrix entries are not easily accessible. The bulk of our method is therefore dedicated to forming approximations of these matrix entries that can be evaluated rapidly.

Lumped mass matrix. At the continuum level, Φ is assumed to be non-negative. However, entries of Φ_h , involve inverse mass matrices, which typically contain negative numbers. If there are too many negative numbers, or if the negative numbers are too large, our algorithm will be less robust. We therefore recommend replacing the mass matrix, \mathbf{M} , with a positive diagonal *lumped mass* approximation. In our numerical results, we use the lumped mass matrix constructed by replacing off-diagonal entries of the mass matrix by zero. Other mass lumping techniques may be used [43].

4. Key innovations. In this section we present two key innovations that our method is based on. First, we define moments of the impulse responses, ϕ_x , show how these moments can be computed efficiently, and use these moments to form ellipsoid shaped a-priori estimates for the supports of the impulse responses (Section 4.1). Second, we describe an improved method of approximating impulse responses by other nearby impulse responses, which we call “normalized local mean displacement invariance” (Section 4.2).

282 **4.1. Impulse response moments and ellipsoid support estimate.** The
 283 impulse response ϕ_x may be interpreted as a scaled probability distribution because
 284 of the non-negative integral kernel property. Let $V : \Omega \rightarrow \mathbb{R}$,

285 (4.1)
$$V(x) := \int_{\Omega} \phi_x(y) dy$$

286 denote the spatially varying scaling factor, and for $i, j = 1, \dots, d$ define $\mu : \Omega \rightarrow \mathbb{R}^d$
 287 and $\Sigma : \Omega \rightarrow \mathbb{R}^{d \times d}$ as follows:

288 (4.2)
$$\mu^i(x) := \int_{\Omega} (\phi_x(y)/V(x)) y^i dy$$

289 (4.3)
$$\Sigma^{ij}(x) := \int_{\Omega} (\phi_x(y)/V(x)) (y^i - \mu^i(x)) (y^j - \mu^j(x)) dy,$$

291 where $\mu^i(x)$ denotes the i^{th} component of the vector $\mu(x)$, and $\Sigma^{ij}(x)$ denotes the
 292 (i, j) entry of the matrix $\Sigma(x)$. The vector $\mu(x) \in \mathbb{R}^d$ and the matrix $\Sigma(x) \in \mathbb{R}^{d \times d}$
 293 are the mean and covariance of the normalized version of ϕ_x , respectively.

294 The direct approach to compute $V(x)$, $\mu(x)$, and $\Sigma(x)$ is to apply \mathcal{A} to a point
 295 source centered at x to obtain ϕ_x , as per (3.3). Then one can post process ϕ_x to
 296 determine $V(x)$, $\mu(x)$, and $\Sigma(x)$. However, this direct approach is not feasible be-
 297 cause we need to know $V(x)$, $\mu(x)$, and $\Sigma(x)$ before we compute ϕ_x , in order choose
 298 the point x . Computing ϕ_x in order to determine $V(x)$, $\mu(x)$, and $\Sigma(x)$ would be
 299 extremely computationally expensive, and defeat the purpose of our algorithm, which
 300 is to reduce the computational cost by computing impulse responses in batches. For-
 301 tunately, it is possible to compute $V(x)$, $\mu(x)$, and $\Sigma(x)$ indirectly, *for all points*
 302 $x \in \Omega$ simultaneously, by applying \mathcal{A}^T to one constant function, d linear functions,
 303 and $d(d+1)/2$ quadratic functions (e.g., 6 total operator applies in two spatial di-
 304 mensions and 10 in three spatial dimensions). This may be motivated by analogy to
 305 matrices. If $\mathbf{B} \in \mathbb{R}^{N \times N}$ is a matrix with i^{th} column \mathbf{b}_i and $\mathbf{w} \in \mathbb{R}^N$, then

306
$$\mathbf{B}^T \mathbf{w} = \begin{bmatrix} - & \mathbf{b}_1^T & - \\ \vdots & & \vdots \\ - & \mathbf{b}_N^T & - \end{bmatrix} \mathbf{w} = \begin{bmatrix} \mathbf{b}_1^T \mathbf{w} \\ \vdots \\ \mathbf{b}_N^T \mathbf{w} \end{bmatrix}.$$

307 By computing one matrix-vector product of \mathbf{B}^T with \mathbf{w} , we compute the inner product
 308 of each column of \mathbf{B} with \mathbf{w} simultaneously. The operator case is analogous, with ϕ_x
 309 taking the place of a matrix column. We have

310 (4.4)
$$(\mathcal{A}^T w)^*(x) = \int_{\Omega} \Phi(y, x) w(y) dy = (\phi_x, w)_{L^2(\Omega)}.$$

311 By computing one operator application of \mathcal{A}^T to w , we compute the inner product of
 312 each ϕ_x with w , for all points x simultaneously.

313 Let C , L^i , and Q^{ij} be the following constant, linear, and quadratic functions:

314
$$C(x) := 1, \quad L^i(x) := x^i, \quad Q^{ij}(x) := x^i x^j$$

315 for $i, j = 1, \dots, d$. Using the definition of V in (4.1) and using (4.4), we have

316
$$V(x) = \int_{\Omega} \phi_x(y) C(y) dy = (\phi_x, C)_{L^2(\Omega)} = (\mathcal{A}^T C)^*(x).$$

317 Hence, we compute $V(x)$ for all x simultaneously by applying \mathcal{A}^T to C . Analogous
 318 manipulations show that $\mu(x)$ and $\Sigma(x)$ may be computed for all points x simultane-
 319 ously by applying \mathcal{A}^T to the functions L^i and Q^{ij} , respectively. We have

320 (4.5a)
$$V = (\mathcal{A}^T C)^*$$

321 (4.5b)
$$\mu^i = (\mathcal{A}^T L^i)^* / V$$

323 (4.5c)
$$\Sigma^{ij} = (\mathcal{A}^T Q^{ij})^* / V - \mu^i \cdot \mu^j$$

324 for $i, j = 1, \dots, d$. Here u/w denotes pointwise division, $(u/w)(x) = u(x)/w(x)$, and
 325 $u \cdot w$ denotes pointwise multiplication, $(u \cdot w)(x) = u(x)w(x)$.

326 We approximate the support of ϕ_x with the ellipsoid

327 (4.6)
$$E_x := \{x' \in \Omega : (x' - \mu(x))^T \Sigma(x)^{-1} (x' - \mu(x)) \leq \tau^2\},$$

328 where τ is a fixed constant. The ellipsoid E_x is the set of points within τ standard
 329 deviations from the mean of the Gaussian distribution with mean $\mu(x)$ and covariance
 330 $\Sigma(x)$, i.e., the Gaussian distribution which has the same mean and covariance as
 331 the normalized version of ϕ_x . The quantity τ is a parameter that must be chosen
 332 appropriately. The larger τ is, the larger the ellipsoid E_x is, and the more conservative
 333 the estimate is for the support of ϕ_x . However, in Section 5.1 we will see that the
 334 cost of our algorithm depends on how many non-overlapping ellipsoids E_x we can
 335 “pack” in the domain Ω (more ellipsoids is better), and choosing a larger value of
 336 τ means that fewer ellipsoids will fit in Ω . In practice, we find that $\tau = 3$ yields a
 337 reasonable balance between these competing interests, and use $\tau = 3$ in our numerical
 338 results. The fraction of the “mass” of ϕ_x residing outside of E_x is less than $1/\tau^2$ by
 339 Chebyshev’s inequality, though this bound is typically conservative.

340 **4.2. Local mean displacement invariance.** Let x and x' be points in Ω that
 341 are “close” to each other, and consider the following approximations:

342 (4.7)
$$\phi_x(y) \approx \phi_{x'}(y)$$

343 (4.8)
$$\phi_x(y) \approx \phi_{x'}(y - x + x')$$

344 (4.9)
$$\phi_x(y) \approx \phi_{x'}(y - \mu(x) + \mu(x'))$$

345 (4.10)
$$\phi_x(y)/V(x) \approx \phi_{x'}(y - \mu(x) + \mu(x'))/V(x').$$

347 These are four different ways to approximate an impulse response by a nearby im-
 348 pulse response, with each successive approximation building upon the previous ones.
 349 Our method uses (4.10), which is the most sophisticated of these approximations.
 350 Approximation (4.7) says that ϕ_x can be approximated by $\phi_{x'}$ when x and x' are
 351 close. Operators satisfying (4.7) can be well approximated via low rank CUR ap-
 352 proximation [54]. However, the required rank in the low rank approximation can be
 353 large, which makes algorithms based on (4.7) expensive. Operators that satisfy (4.8)
 354 are called “locally translation invariant” because integral kernel entries $\Phi(y, x)$ for
 355 such operators are approximately invariant under translation of x and y by the same
 356 displacement, i.e., $x \rightarrow x + h$ and $y \rightarrow y + h$. It is straightforward to show that if
 357 equality holds in (4.8), then \mathcal{A} is a convolution operator. Locally translation invariant
 358 operators act like convolutions locally, and can therefore be well approximated by PC
 359 approximations.

360 Approximation (4.9) improves upon (4.7) and (4.8), and generalizes both. On
 361 one hand, if (4.7) holds, then $\mu(x) \approx \mu(x')$, and so (4.9) holds. On the other hand,

362 translating a distribution translates its mean, so if (4.8) holds, then $\mu(x') - \mu(x) \approx x' -$
 363 x , so again (4.9) holds. But approximation (4.9) can hold in situations where neither
 364 (4.7) nor (4.8) holds. For example, because the expected value commutes with affine
 365 transformations, (4.9) will hold when \mathcal{A} is locally translation invariant with respect to
 366 a rotated frame of reference, while (4.8) will not hold in this case. Additionally, (4.9)
 367 generalizes to operators $\mathcal{A} : L^2(\Omega_1) \rightarrow L^2(\Omega_2)'$ that map between function spaces
 368 on different domains Ω_1 and Ω_2 , and even operators that map between domains
 369 with different spatial dimensions. In contrast, (4.8) does not naturally generalize to
 370 operators that map between function spaces on different domains, because the formula
 371 $y - x + x'$ requires vectors in Ω_2 and Ω_1 to be added together. We call (4.9) “local
 372 mean displacement invariance,” and illustrate (4.9) in Figure 2.

373 We use approximation (4.10), which is the same as (4.9), except for the extra
 374 factors of $1/V(x)$ on the left hand side and $1/V(x')$ on the right hand side. These
 375 factors make the approximation more robust if $V(x)$ varies widely. Approximation
 376 (4.10) is equivalent to (4.9), but with ϕ_x replaced by its normalized version, $\phi_x/V(x)$.
 377 We call (4.10) “normalized local mean displacement invariance.”

378 **5. Operator approximation algorithm.** We use (4.5) in Section 4.1 to com-
 379 pute V , μ , and Σ by applying \mathcal{A}^T to a small number of polynomial functions. Then
 380 we use (4.6) to form ellipsoid shaped estimates for the support of the ϕ_x , *without*
 381 computing them. This allows us to compute large numbers of ϕ_{x_i} in “batches,” η_b
 382 (see Figure 1). We compute one batch, denoted η_b , by applying \mathcal{A} to a weighted
 383 sum of point sources (Dirac comb) associated with a batch, S_b , of points x_i scattered
 384 throughout Ω (Section 5.2). The batch of points S_b is chosen via a greedy ellipsoid
 385 packing algorithm so that, for $x_i, x_j \in S_b$, the support ellipsoid for ϕ_{x_i} and the sup-
 386 port ellipsoid for ϕ_{x_j} do not overlap if $i \neq j$ (Section 5.1). Because these supports
 387 do not overlap (or do not overlap much), we can post process η_b to recover the func-
 388 tions ϕ_{x_i} associated with all points $x_i \in S_b$ —with one application of \mathcal{A} , we recover
 389 many ϕ_{x_i} (Section 5.2). The process is repeated until a desired number of batches is
 390 reached.

391 Once the batches η_b are computed, we approximate the integral kernel $\Phi(y, x)$
 392 at arbitrary points (y, x) by interpolation of translated and scaled versions of the
 393 computed ϕ_{x_i} (Section 5.3). The key idea behind the interpolation is the normalized
 394 local mean displacement invariance assumption discussed in Section 4.2. Specifi-
 395 cally, we approximate $\Phi(y, x) = \phi_x(y)$ by a weighted linear combination of the values
 396 $\frac{V(x)}{V(x_i)}\phi_{x_i}(y - \mu(x) + \mu(x_i))$ for a small number of sample points x_i near x . The weights
 397 are determined by radial basis function interpolation.

398 The ability to rapidly evaluate approximate kernel entries $\Phi(y, x)$ allows us to
 399 construct an H-matrix approximation, $\Phi_H \approx \Phi$, using the conventional adaptive
 400 cross H-matrix construction method (Section 5.4). In this method, one forms low rank
 401 approximations of off-diagonal blocks of the matrix by sampling rows and columns
 402 of those blocks. We then use H-matrix methods to convert Φ_H into an H-matrix
 403 approximation $\mathbf{A}_H \approx \mathbf{A}$. The complete algorithm for constructing \mathbf{A}_H is shown in
 404 Algorithm 1. The cost is discussed in Section 6.

405 When \mathcal{A} is symmetric positive semi-definite, \mathbf{A}_H may be non-symmetric and
 406 indefinite due to errors in the approximation. In this case, one may (optionally)
 407 modify the H-matrix representation of \mathbf{A}_H to make it symmetric positive semi-definite
 408 using a rational function method that we describe in Appendix B.

Algorithm 1: Construct H-matrix approximation

Input : Linear operator \mathcal{A} , parameter n_b
Output: H-matrix \mathbf{A}_H

- 1 Compute V, μ , and Σ (Equations (4.5) in Section 4.1)
- 2 **for** $k = 1, 2, \dots, n_b$ **do**
- 3 Choose a batch of sample points, S_k (Section 5.1)
- 4 Compute η_k by applying \mathcal{A} to the Dirac comb for S_k (Section 5.2)
- 5 Form H-matrix approximation Φ_H of integral kernel (Section 5.4)
- 6 Form H-matrix approximation \mathbf{A}_H of \mathcal{A} (Section 5.4)
- 7 (optional) Modify \mathbf{A}_H to make it symmetric positive definite (Appendix B)

409 **5.1. Sample point selection via greedy ellipsoid packing.** We choose sam-
410 ple points, x_i , in batches S_k . We use a greedy ellipsoid packing algorithm to choose
411 as many points as possible per batch, while ensuring that there is no overlap between
412 the support ellipsoids, E_{x_i} , associated with the sample points within a batch.

413 We start with a finite set of candidate points X and build S_k incrementally with
414 points selected from X . For simplicity of explanation, here S_k and X are mutable
415 sets that we add points to and remove points from. First we initialize S_k as an empty
416 set. Then we select the candidate point $x_i \in X$ that is the farthest away from all
417 points in previous sample point batches $S_1 \cup S_2 \cup \dots \cup S_{k-1}$. Candidate points for the
418 first batch S_1 are chosen randomly from X . Once x_i is selected, we remove x_i from
419 X . Then we perform the following two checks:

- 420 1. We check whether x_i is sufficiently far from all of the previously chosen points
421 in the current batch, in the sense that $E_{x_i} \cap E_{x_j} = \emptyset$ for all $x_j \in S_k$.
422 2. We make sure that $V(x_i)$ is not too small, by checking whether $V(x_i) > \epsilon V_{\max}$.
423 Here V_{\max} is the largest value of $V(x_j)$ over all points q in the initial set of
424 candidate points, and ϵ is a small threshold parameter (we use $\epsilon = 10^{-5}$).

425 If x_i passes both these checks (i.e., if x_i is sufficiently far from other points in the
426 batch, and $V(x_i)$ is not too small) then we add x_i to S_k . Otherwise we discard x_i . This
427 process repeats until there are no more points in X . We check whether $E_{x_i} \cap E_{x_j} = \emptyset$
428 using the ellipsoid intersection test described in Appendix A. We repeat the process
429 to construct several batches of points S_1, S_2, \dots, S_{n_b} . In our implementation, for
430 each batch, X is initialized as the set of all Lagrange nodes for the finite element
431 basis functions used to discretize the problem, except for points in previously chosen
432 batches.

433 **5.2. Impulse response batches.** We compute impulse responses ϕ_{x_i} in batches
434 by applying \mathcal{A} to Dirac combs. The Dirac comb, ξ_k , associated with a batch of
435 sample points, S_k , is the following weighted sum of Dirac distributions (point sources)
436 centered at the points $x_i \in S_k$:

437
$$\xi_k := \sum_{x_i \in S_k} \delta_{x_i} / V(x_i).$$

438 We compute the *impulse batch*, η_k , by applying \mathcal{A} to the Dirac comb:

439 (5.1)
$$\eta_k := (\mathcal{A}\xi_k^*)^* = \sum_{x_i \in S_k} \phi_{x_i} / V(x_i).$$

440 The last equality in (5.1) follows from linearity and the definition of ϕ_{x_i} in (3.3). Since
 441 the points x_i are chosen so that the ellipsoid E_{x_i} that (approximately) supports ϕ_i ,
 442 and the ellipsoid E_{x_j} that (approximately) supports ϕ_j do not overlap when $i \neq j$,
 443 we have (approximately)

$$444 \quad (5.2) \quad \phi_{x_i}(z) = \begin{cases} \eta_k(z)V(x_i), & z \in E_{x_i} \\ 0, & \text{otherwise} \end{cases}$$

445 for all $x_i \in S_k$. By performing one matrix-vector product, $\xi_k \mapsto (\mathcal{A}\xi_k^*)^*$, we recover
 446 ϕ_{x_i} for every point $x_i \in S_k$.

447 Each point source δ_{x_i} is scaled by $1/V(x_i)$ so that the resulting scaled impulse
 448 responses within η_k are comparable in magnitude. Without this scaling, the portion
 449 of ϕ_{x_i} outside of E_{x_i} , which we neglect, may overwhelm ϕ_{x_j} for a nearby point x_j
 450 if $V(x_i)$ is much larger than $V(x_j)$. Note that we are not in danger of dividing
 451 by zero, because our ellipsoid packing procedure from Section 5.1 excludes x_i from
 452 consideration as a sample point if $V(x_i)$ is smaller than a predetermined threshold.

453 **5.3. Approximate integral kernel entries.** Given $(y, x) \in \Omega \times \Omega$, let

$$454 \quad z_i := y - \mu(x) + \mu(x_i), \quad i = 1, \dots, k_n,$$

455 where $\{x_i\}_{i=1}^{k_n}$ are the k_n nearest sample points to x , excluding points x_i for which
 456 $z_i \notin \Omega$. Here k_n is a small user defined parameter; e.g., $k_n = 10$. We find the k_n
 457 nearest sample points to x by querying a precomputed kd-tree [9] of all sample points.
 458 We check whether $z_i \in \Omega$ by querying a precomputed axis aligned bounding box tree
 459 (AABB tree) [21] of the mesh cells used to discretize the problem. Further, let

$$460 \quad (5.3) \quad f_i := \frac{V(x)}{V(x_i)} \phi_{x_i}(z_i), \quad i = 1, \dots, k_n.$$

461 Note that $\phi_{x_i}(z_i)$ is well defined because $z_i \in \Omega$, and $\frac{V(x)}{V(x_i)}$ is well defined because
 462 our sample point picking procedure in Section 5.1 ensures that $V(x_i) > 0$. Per the
 463 discussion in Section 4.2, we expect $\Phi(y, x) \approx f_i$ for $i = 1, \dots, k_n$; the closer x_i is to
 464 x , the better we expect the approximation to be. We therefore approximate $\Phi(y, x)$
 465 by interpolating the (point,value) pairs $\{(x_i, f_i)\}_{i=1}^{k_n}$ at the point x . Interpolation is
 466 performed using the following radial basis function [71] scheme:

$$467 \quad (5.4) \quad \Phi(y, x) \approx \tilde{\Phi}(y, x) := \sum_{i=1}^{k_n} c_i \varphi(\|x - x_i\|),$$

468 where c_i are weights, and φ is a polyharmonic spline radial basis function, which is
 469 defined as $\varphi(r) := r^d$ if d is odd, and $\varphi(r) := r^d \log r$ for $r > 0$ and $\varphi(0) := 0$ if d
 470 is even. The vector of weights, $c = (c_1, c_2, \dots, c_{k_n})^T$, is found as the solution to the
 471 $k_n \times k_n$ linear system

$$472 \quad (5.5) \quad Bc = f,$$

473 where $B \in \mathbb{R}^{k_n \times k_n}$ has entries $B_{ij} := \varphi(\|x_i - x_j\|)$, and $f \in \mathbb{R}^{k_n}$ has entries f_i from
 474 (5.3). Polyharmonic spline radial basis function interpolation yields the smoothest
 475 interpolant of the given points and values, in a certain least squares sense [19].

Symbol	Typical size	Variable name
N	$10^2\text{--}10^9$	Number of finite element degrees of freedom
n_b	1–25	Number of batches
k_h	5–50	H-matrix rank
k_n	5–15	Number of nearest neighbors for RBF interpolation
d	1–3	Spatial dimension
m	$10^1\text{--}10^4$	Total number of sample points (all batches)
$ S_i $	1–500	Number of sample points in the i th batch
l	1–2	Rational function parameter for SPSD modification

Table 1: Symbols used for variables in computational cost estimates, and approximate ranges for their sizes in practice.

476 To evaluate f_i , we check whether $z_i \in E_{x_i}$ using (4.6). If $z_i \notin E_{x_i}$, then z_i is
 477 outside the estimated support of ϕ_{x_i} , so we set $f_i = 0$. If $z_i \in E_{x_i}$, we look up the
 478 batch index b such that $x_i \in S_b$, and evaluate f_i via the formula $f_i = V(x)\eta_b(z_i)$,
 479 per (5.2). Note that z_i is typically not a gridpoint of the mesh used to discretize
 480 the problem, even if y , x , and x_i are gridpoints. Hence, evaluating $\eta_b(z_i)$ requires
 481 determining which mesh cell contains z_i , then evaluating finite element basis functions
 482 on that mesh cell. Fortunately, the mesh cell containing z_i was determined as a side
 483 effect of querying the AABB tree of mesh cells when we checked whether $z_i \in \Omega$.

484 **5.4. Hierarchical matrix construction.** We form our H-matrix approxima-
 485 tion $\mathbf{A}_H \approx \mathbf{A}$ by forming H-matrix representations Φ_H and \mathbf{M}_H of Φ and \mathbf{M} , re-
 486 spectively, then using fast H-matrix methods to multiply these matrices per (3.5) to
 487 form $\mathbf{A}_H = \mathbf{M}_H \Phi_H \mathbf{M}_H$. We use H1 matrices in our numerical results, but any other
 488 H-matrix format could be used instead. For more details on H-matrices, see [39].

489 We form Φ_H using the standard geometrical clustering/adaptive cross method im-
 490 plemented within the HLIBpro software package [48]. For details about the algorithms
 491 used for geometrical clustering, H-matrix construction, and H-matrix operations in
 492 HLIBpro, we recommend the the following papers [11, 36, 49]. Although Φ is a dense
 493 $N \times N$ matrix, constructing Φ_H only requires evaluation of $O(k_h^2 N \log N)$ kernel en-
 494 tries $\Phi_{ij} = \tilde{\Phi}(\zeta_i, \zeta_j)$, and these entries are computed via the radial basis function
 495 interpolation method described in Section 5.3. Here k_h is the rank of the highest rank
 496 block in the H-matrix. A dense representation of Φ is never formed. We form \mathbf{M}_H
 497 using standard H-matrix methods for sparse matrices implemented within HLIBpro,
 498 using the same geometrical clustering block cluster tree as was used for Φ .

499 **6. Computational cost.** The computational cost of our method may be divided
 500 into the costs to perform the following tasks: (1) Computing impulse moments and
 501 batches; (2) Building the H-matrix; (3) Performing linear algebra operations with the
 502 H-matrix. This may optionally include modifying the H-matrix to make it symmetric
 503 positive semi-definite. In practice, (1) is the dominant cost, because (1) is the only
 504 task that requires applying \mathcal{A} to vectors, and our method is targeted at applications
 505 in which applying \mathcal{A} to a vector requires an expensive computational procedure such
 506 as solving a large scale PDE. All operations that do not require applications of \mathcal{A} to
 507 vectors are nearly linear, and therefore scalable, in the size of the problem, N . We
 508 now describe these costs in detail. For convenience, Table 1 lists variable symbols and
 509 their approximate sizes.

510 (1) *Computing impulse response moments and batches.* Computing V , μ , and Σ
 511 requires applying \mathcal{A} to 1, d , and $d(d+1)/2$ vectors, respectively. Computing each η_i
 512 requires applying \mathcal{A} to one vector, so computing $\{\eta_i\}_{i=1}^{n_b}$ requires n_b operator applies.
 513 In total, computing all impulse response moments and batches therefore requires

$$514 \quad 1 + d + d(d+1)/2 + n_b \quad \text{operator applies.}$$

515 In a typical application one might have $d = 2$ and $n_b = 5$, in which case a modest 11
 516 operator applies are required.

517 Computing the impulse response batches also requires choosing sample point
 518 batches via the greedy ellipsoid packing algorithm described in Section 5.1. Choosing
 519 the i th batch of sample points may require performing up to $N|S_i|$ ellipsoid intersec-
 520 tion tests, where $|S_i|$ is the number of sample points in the i th batch. Choosing all
 521 of the sample points therefore requires performing at most

$$522 \quad Nm \quad \text{ellipsoid intersection tests,}$$

523 where m is the total number of sample points in all batches. The multiplicative
 524 dependence of N with m is undesirable since m may be large, and reducing this cost
 525 is a target for future work. However, from a practical perspective, the cost for picking
 526 sample points is small compared to other parts of the approximation algorithm, so
 527 we do not pursue these improvements in this paper.

528 (2) *Building the H-matrix.* The classical H-matrix construction technique requires
 529 evaluating $O(k_h^2 N \log N)$ matrix entries of the approximation, where k_h is the H-
 530 matrix rank, i.e., the maximum rank among the blocks of the H-matrix. To evaluate
 531 one matrix entry, first one must find the k_n nearest sample points to a given point,
 532 where k_n is the number of impulse responses used in the RBF interpolation. This is
 533 done using a precomputed kd-tree of sample points, and requires $O(k_n \log m)$ float-
 534 ing point and logical elementary operations. Second, one must find the mesh cells
 535 that the points $\{z_i\}_{i=1}^{k_n}$ reside. This is done using an AABB tree of mesh cells, and
 536 requires $O(k_n \log N)$ elementary operations. Third, one must evaluate finite element
 537 basis functions on each of those cells, which requires $O(k_n)$ elementary operations.
 538 Finally, one must solve a $k_n \times k_n$ linear system to perform the radial basis function
 539 interpolation, which requires $O(k_n^3)$ elementary operations, which typically is negligible.
 540 Therefore, building the H-matrix requires

$$541 \quad O((k_h^2 N \log N)(k_n \log N + k_n^3)) \quad \text{elementary operations.}$$

542 (3) *Performing linear algebra operations with the H-matrix.* It is well known that
 543 H-matrix methods for matrix-vector products, matrix-matrix addition, matrix-matrix
 544 multiplication, matrix factorization, and matrix inversion, and low rank updates re-
 545 quire

$$546 \quad (6.1) \quad O(k_h^2 N \log(N)^a) \quad \text{elementary operations,}$$

547 where $a \in \{0, 1, 2, 3\}$ [35]. The precise cost depends on the type of H-matrix used
 548 and the operation being performed. In the numerical results (Section 7), we use one
 549 matrix-matrix addition to add the H-matrix approximation of the data misfit term
 550 in the Hessian to the regularization term in the Hessian. Constructing an H-matrix
 551 representation of the regularization term in the Hessian requires one matrix inversion
 552 and two matrix-matrix multiplications. Then, we use one matrix factorization to form

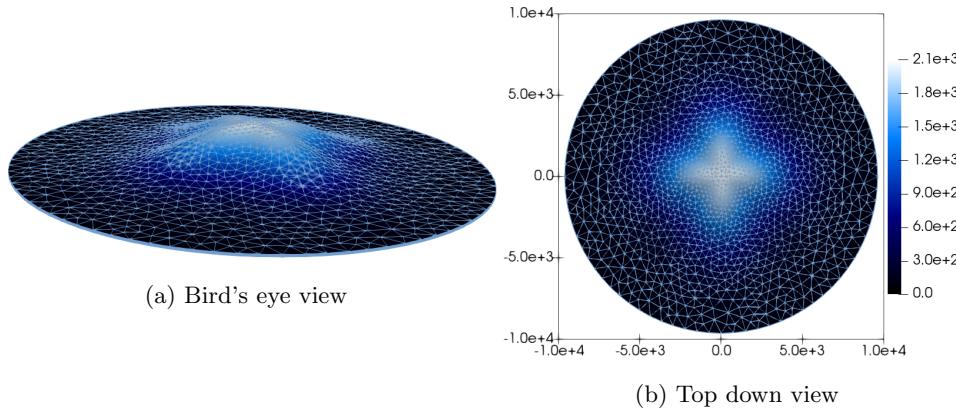


Fig. 3: The ice sheet, discretized by a mesh of tetrahedra. Color indicates the height of the base of the ice sheet (i.e., the mountain topography). The radius of the domain is 10^4 meters, the maximum height of the mountain is 2.1×10^3 meters, and the average thickness of the ice sheet is 250 meters.

553 a preconditioner for the Newton linear system. The rational method for modifying \mathbf{A}_H
 554 to be positive semi-definite, described in Appendix B, is performed using two matrix-
 555 matrix additions, $l+1$ matrix-matrix multiplications, and one matrix inversion, where
 556 2^l is the rational function order.

557 7. Numerical results.

558 7.1. Example 1: Inversion for the basal sliding coefficient in an ice
 559 sheet flow problem. We use our method to approximate the data misfit Gauss-
 560 Newton Hessian in an inverse problem governed by a PDE which models steady state
 561 ice sheet flow [61]. Briefly, a sheet of ice is flowing down a mountain (see Figure
 562 3). Given observations of the tangential component of the ice velocity on the top
 563 surface of the ice, we invert for the logarithm of the unknown spatially varying basal
 564 sliding parameter field, which governs the resistance to sliding along the base of the
 565 ice sheet. The inverse problem considered here is synthetic, but in practice surface
 566 velocity observations are available from satellites.

The domain that is filled with ice is denoted by $U \subset \mathbb{R}^3$. The basal, lateral and top parts of ∂U are denoted by Γ_b , Γ_l , and Γ_t , respectively. The governing equations are the linear incompressible Stokes equations,

570 (7.1a) $-\nabla \cdot \sigma = f$ in U ,
 571 (7.1b) $\nabla \cdot v = 0$ in U ,
 572 (7.1c) $\sigma \nu = 0$ on Γ_t ,
 573 (7.1d) $v \cdot \nu = 0$ and $T(\sigma \nu + \exp(q)v) = 0$ on Γ_b ,
 574 (7.1e) $\sigma \nu + sv = 0$ on Γ_l .

The solution to these equations is the pair (v, p) , where v is the ice flow velocity field, and p is the pressure field. Here, q is the unknown logarithmic basal sliding field

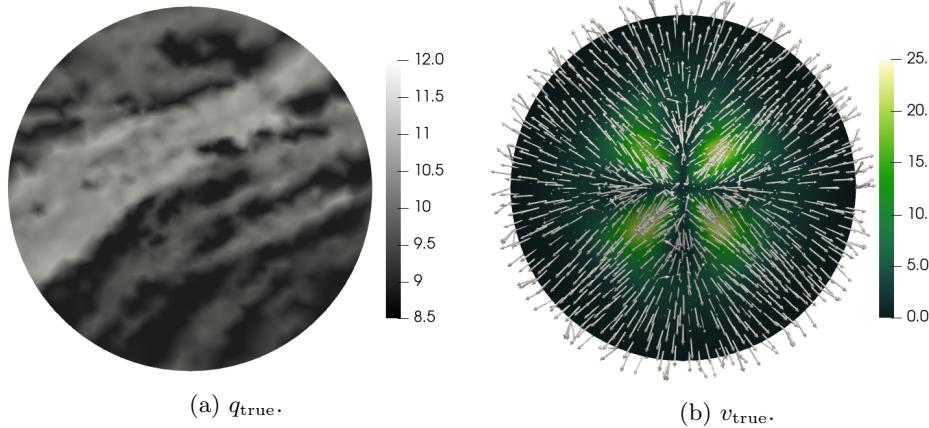


Fig. 4: True parameter, q_{true} , (left) and true velocity v_{true} (right). In the right plot, arrows indicate the direction of v_{true} and color indicates the magnitude of v_{true} .

(note that large q corresponds to high resistance to sliding) which is defined on the 2D surface Γ_b . The quantity f is the body force density due to gravity, $s = 10^6$ is a Robin boundary condition constant, ν is the outward normal and T is the tangential projection operator that restricts a vector field to its tangential component along the boundary. We employ Glen's flow law [34], $\sigma = 2\eta\dot{\varepsilon} - Ip$, which is a constitutive law for ice that relates the stress tensor, σ , to the strain rate tensor, $\dot{\varepsilon} = \frac{1}{2}(\nabla v + \nabla v^\top)$. Here η is the effective viscosity and I is the unit matrix. Glen's exponent has been chosen to be one, which provides for a linear Stokes model. Note that while the PDE is linear, the parameter-to-solution map, $q \mapsto (v, p)$, is nonlinear.

The pressure, p , is discretized with first order scalar continuous Galerkin finite elements defined on a mesh of tetrahedra. The velocity, v , is discretized with second order continuous Galerkin finite elements on the same mesh. The parameter q is discretized with first order scalar continuous Galerkin finite elements on the mesh of triangles that results from restricting the tetrahedral mesh to the basal boundary, Γ_b . Note that Γ_b is a 2D surface embedded in 3D due to the mountain topography. We also generate a flattened version of Γ_b , denoted by $\Omega \subset \mathbb{R}^2$, by ignoring the height coordinate. The parameter q is dually viewed as a function on Γ_b for the purpose of solving the Stokes equations, and as a function on Ω for the purpose of building our Hessian approximations and defining the regularization. The observations are generated by adding 1% multiplicative Gaussian noise to the tangential component of the velocity field restricted to the top surface of the geometry. The true basal sliding coefficient field and the true velocity field, which is obtained by solving (7.1), are shown in Figure 4.

To reconstruct the basal sliding parameter field q , we formulate an inverse problem governed by the Stokes problem (7.1). In particular this problem is formulated as a nonlinear least squares optimization problem, whose cost function consists of a misfit term (between the observations and model output) and a bi-Laplacian regularization term following [45]. The regularization is centered at the constant function $q_0(x) = 10.5$. To mitigate boundary effects we use a constant coefficient Robin boundary condition as in [62]. The parameters are chosen so that the Green's function of

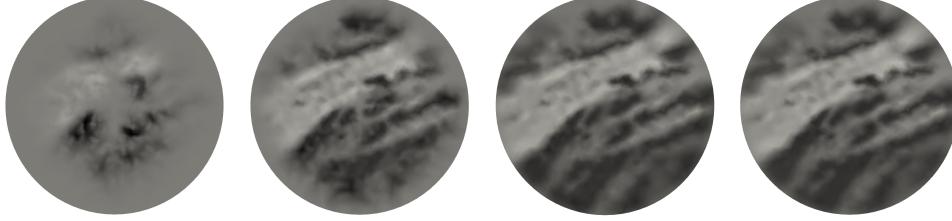


Fig. 5: The log basal sliding parameter after 1 (left), 3 (second from left), and 4 (third from left) Newton iterations, and the final Newton iterate (right). The PSF (5) preconditioner is constructed between Newton iterations 3 and 4.

608 the Hessian of the regularization has a characteristic length of 0.25 of the domain
 609 radius. For the specific setup, we refer the reader to [69, Section 2.2]. We choose the
 610 regularization parameter, γ , that controls the overall strength of the regularization
 611 using Morozov's discrepancy principle. The optimal value of γ is 7.3×10^3 , which is
 612 used for all the numerical results except for Figure 6b, where we vary γ to study how
 613 changing γ impacts the effectiveness of the proposed preconditioner.

614 We solve the inverse problem with an inexact Newton preconditioned conjugate
 615 gradient (PCG) scheme and a globalizing Armijo line search [59]. The (discrete)
 616 Newton updates, $\hat{\mathbf{q}}_k$, are obtained by solving

$$617 \quad (7.2) \quad \mathbf{H}\hat{\mathbf{q}}_k = -\mathbf{g}_k \quad \text{or} \quad \mathbf{H}_{\text{gn}}\hat{\mathbf{q}}_k = -\mathbf{g}_k,$$

618 wherein we choose the initial guess as the discretization of the constant function q_0 .
 619 Here \mathbf{g}_k , \mathbf{H} and \mathbf{H}_{gn} are the discretized gradient, Hessian, and Gauss-Newton Hessian
 620 of the inverse problem cost function, respectively. To ensure positive definiteness of
 621 the Hessian we use \mathbf{H}_{gn} for the first five iterations, and \mathbf{H} for all subsequent iterations.
 622 The Newton iterations are terminated when $\|\mathbf{g}_k\|/\|\mathbf{g}_0\| < 10^{-8}$. Systems (7.2) are
 623 solved inexactly using an inner PCG iteration, which terminates when the norm of
 624 the linear system residual is less than $\min(0.5, \sqrt{\|\mathbf{g}_k\|}) \|\mathbf{g}_k\|$.

625 We use the framework described in this paper to generate Hessian precondition-
 626 ers. We build H-matrix approximations of the data misfit Gauss-Newton Hessian (the
 627 term in \mathbf{H}_{gn} that arises from the data misfit) using 1, 5, and 25 batches. These approx-
 628 imations are denoted by PSF (1), PSF (5), and PSF (25), respectively. The resulting
 629 H-matrices, which may be indefinite due to approximation error, are modified to be
 630 symmetric positive definite via the procedure described in Appendix B, with $l = 2$.
 631 We approximate \mathbf{H}_{gn} rather than \mathbf{H} because \mathbf{H} more often has negative values in its
 632 integral kernel, especially when one is away from the optimal basal sliding parameter.
 633 Our results show that our preconditioners are also good preconditioners for \mathbf{H} . The
 634 H-matrix approximation of the data misfit Gauss-Newton Hessian is added to an H-
 635 matrix approximation of the Hessian of the regularization term to form an H-matrix
 636 approximation of \mathbf{H}_{gn} . We note that the Hessian of the regularization is a combina-
 637 tion of stiffness and mass matrices, which are sparse. Therefore, we form H-matrix
 638 representations of these matrices and combine them using standard sparse H-matrix
 639 techniques. We factor the overall Gauss-Newton Hessian approximation, denoted $\widetilde{\mathbf{H}}$,
 640 using fast H-matrix methods, then use the factorization as a preconditioner.

641 Table 2 shows the performance of our preconditioner for accelerating the solution
 642 of optimization problem (??). We build the PSF (5) preconditioner in the third

Iter	PSF (5)			REG			NONE		
	#CG	#Stokes	$\ \mathbf{g}\ $	#CG	#Stokes	$\ \mathbf{g}\ $	#CG	#Stokes	$\ \mathbf{g}\ $
0	1	4	2.0e+7	1	4	2.0e+7	1	4	2.0e+7
1	2	6	5.8e+6	2	6	7.9e+6	2	6	5.8e+6
2	4	10	2.4e+6	5	12	3.9e+6	4	10	2.4e+6
3	2	6+22	5.8e+5	13	28	1.6e+6	13	28	5.8e+5
4	5	12	5.6e+4	42	86	4.8e+5	33	68	9.1e+4
5	10	22	3.5e+3	84	170	7.7e+4	57	116	6.2e+3
6	14	30	2.7e+1	125	252	5.0e+3	76	154	1.1e+2
7	0	2	4.1e-2	194	390	7.9e+1	117	236	2.3e-1
8	—	—	—	0	2	1.6e-1	0	2	4.1e-2
Total	38	114	—	466	950	—	303	624	—

Table 2: Convergence history for solving the Stokes inverse problem using preconditioned inexact Newton PCG. Preconditioners shown are our method with five batches (PSF (5)) constructed at the third iteration, regularization preconditioning (REG), and no preconditioning (NONE). Columns titled #CG show the number of PCG iterations used to solve the Newton system for $\hat{\mathbf{q}}_k$. Columns titled $\|\mathbf{g}\|$ show the l^2 norm of the gradient at \mathbf{q}_k . Columns titled #Stokes show the total number of Stokes PDE solves performed in each Newton iteration. This consists of Stokes solves for u performed during the linesearch going from \mathbf{q}_{k-1} to \mathbf{q}_k , plus one adjoint Stokes solve to compute the gradient at \mathbf{q}_k , plus one incremental forward and one incremental adjoint Stokes solve per PCG iteration for solving the Newton system. In the PSF (5) portion of row 3, we write 6 + 22 to indicate that 6 Stokes solves were used during the standard course of the iteration, and 22 Stokes solves were used to build the PSF (5) preconditioner.

643 Gauss-Newton iteration, and reuse it for all subsequent Gauss-Newton and Newton
 644 iterations. No preconditioning is used in the Gauss-Newton iterations before the PSF
 645 (5) preconditioner is built. We compare our method with the most commonly used
 646 existing preconditioners: no preconditioning (NONE), and preconditioning by the
 647 regularization term in the Hessian (REG). Using PSF (5) reduces the total number of
 648 Stokes PDE solves by a factor of six compared to no preconditioning, and by a factor of
 649 nine compared to regularization preconditioning⁵. In Figure 5 we show select Newton
 650 iterates using our PSF (5) preconditioner.

651 Next, we build PSF (1), PSF (5), and PSF (25) preconditioners based on the
 652 Gauss-Newton Hessian evaluated at the converged solution \mathbf{q} . We use PCG to solve
 653 a linear system with the Hessian as the coefficient operator and a right hand vector
 654 with random i.i.d. entries drawn from the standard normal distribution. In Figure
 655 6a we compare the convergence of PCG for solving this linear system using the PSF
 656 (1), PSF (5), PSF (25), REG, and NONE preconditioners. PCG converges fastest
 657 with the PSF preconditioners, with PSF (25) converging fastest, followed by PSF
 658 (5), followed by PSF (1), as expected. PCG converges much slower with no precon-
 659 ditioning and regularization preconditioning than it does with PSF preconditioning,
 660 with no preconditioning outperforming regularization preconditioning. In Figure 6b,

⁵Interestingly, we observe that no preconditioning outperforms regularization preconditioning here because the noise level (and hence γ) is small.

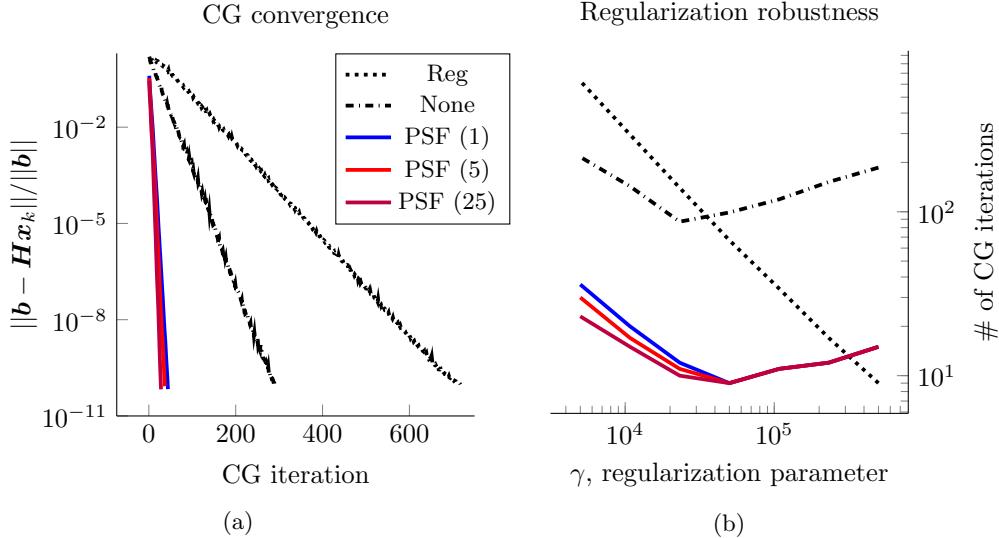


Fig. 6: Left (6a) shows the convergence history for solving $\mathbf{H}\mathbf{x} = \mathbf{b}$ using PCG, where \mathbf{b} has i.i.d. random entries drawn from the standard normal distribution. Here we use $\gamma = 7.3 \times 10^3$, which is determined by the Morozov discrepancy principle. Right (6b) shows the number of PCG iterations required to achieve $\|\mathbf{b} - \mathbf{H}\mathbf{x}_k\|/\|\mathbf{b}\| < 10^{-6}$, where \mathbf{x}_k is the k th PCG iterate, for a range of different γ . Results in these figures are shown for our preconditioner with 1, 5, and 25 batches (PSF (1), PSF (5), and PSF(25), respectively), regularization preconditioning (REG), and no preconditioning (NONE). The preconditioner is constructed using \mathbf{H}_{gn} . For a fair comparison, even as γ changes, \mathbf{H} and \mathbf{H}_{gn} are always evaluated at the same \mathbf{q} , which is the optimal point for the inverse problem with $\gamma = 7.3 \times 10^3$.

661 we perform PCG solves on linear systems of the same form, except now we vary γ .
 662 The performance of our PSF preconditioners is good and relatively stable over a wide
 663 range of γ 's. All PSF preconditioners perform the same for medium and large values
 664 of γ . For small values of γ PSF (25) performs slightly better than PSF (5), which
 665 performs slightly better than PSF (1). As expected, regularization preconditioning
 666 performs well for large γ and poorly for small γ . Our PSF preconditioners outperform
 667 no preconditioning and regularization preconditioning for medium and small γ , and
 668 perform similarly to regularization preconditioning for large γ .

669 In Figure 7a we show the generalized eigenvalues for the generalized eigenvalue
 670 problem $\mathbf{H}\mathbf{u} = \lambda \widetilde{\mathbf{H}}\mathbf{u}$. Here \mathbf{H} is the Hessian evaluated at the reconstructed param-
 671 eter \mathbf{q} for γ chosen to satisfy the Morozov discrepancy principle. The matrix $\widetilde{\mathbf{H}}$ is
 672 one of the PSF (1), PSF (5), or PSF (25) Gauss-Newton Hessian approximations
 673 constructed at that \mathbf{q} , or the regularization Hessian. The PSF preconditioners cluster
 674 the eigenvalues of the Hessian near one, with more batches yielding better clustering.
 675 The regularization preconditioner clusters the trailing eigenvalues but amplifies lead-
 676 ing eigenvalues. In Figure 7b we show the condition numbers of $\widetilde{\mathbf{H}}^{-1}\mathbf{H}$ and $\widetilde{\mathbf{H}}^{-1}\mathbf{H}_{gn}$.
 677 The PSF preconditioners yield the the smallest condition numbers, with more batches
 678 yielding smaller condition numbers. Regularization preconditioning and no precondi-

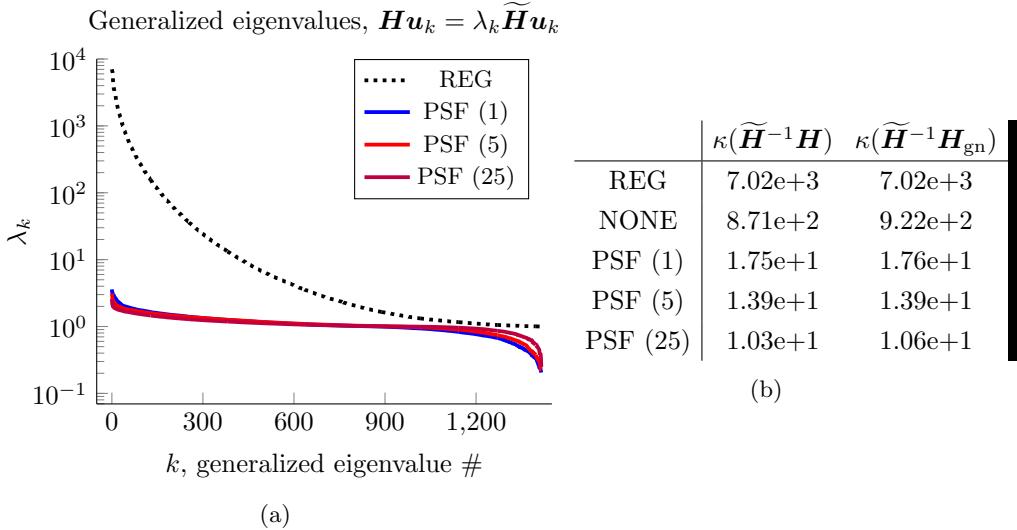


Fig. 7: Left (7a): eigenvalues for generalized eigenvalue problem $\mathbf{H}\mathbf{u}_k = \lambda_k \widetilde{\mathbf{H}}\mathbf{u}_k$. Here \mathbf{H} is the Hessian and $\widetilde{\mathbf{H}}$ is the PSF approximation constructed using the Gauss-Newton Hessian, \mathbf{H}_{gn} , for 1, 5, and 25 batches (PSF (1), PSF (5), and PSF (25), respectively), or the regularization Hessian (REG). Right (7b): condition number for $\widetilde{\mathbf{H}}^{-1}\mathbf{H}$ and $\widetilde{\mathbf{H}}^{-1}\mathbf{H}_{\text{gn}}$ for these different preconditioners, and no preconditioner (NONE). All operators are evaluated at the \mathbf{q} that solves the inverse problem for $\gamma = 7.3 \times 10^3$, as determined by the Morozov discrepancy principle. COND(H)
LABELS!

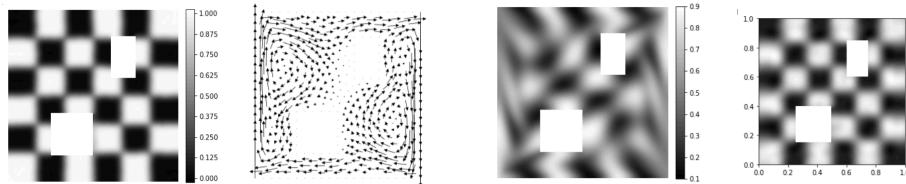


Fig. 8: (a) true initial condition. (b) velocity field. (c) observed concentration at final time. (d) reconstruction of initial condition. black=0, white=1

679 tationing yield larger condition numbers. The condition numbers are similar for both
 680 $\widetilde{\mathbf{H}}^{-1}\mathbf{H}$ and $\widetilde{\mathbf{H}}^{-1}\mathbf{H}_{\text{gn}}$, demonstrating that the preconditioner built based on \mathbf{H}_{gn} is
 681 a good preconditioner for \mathbf{H} .

682 **7.2. Example 2: Inversion for the initial condition in an advection-
 683 diffusion problem.**

- 684 • Picture of velocity field, initial condition, final condition, reconstruction (1)
 685 (OKOK)
 686 • Impulse responses: vary time and Peclet number (3) (OKOK)
 687 • CG convergence (1) CG iters vs Peclet number (Morozov) (2) (OKOK)

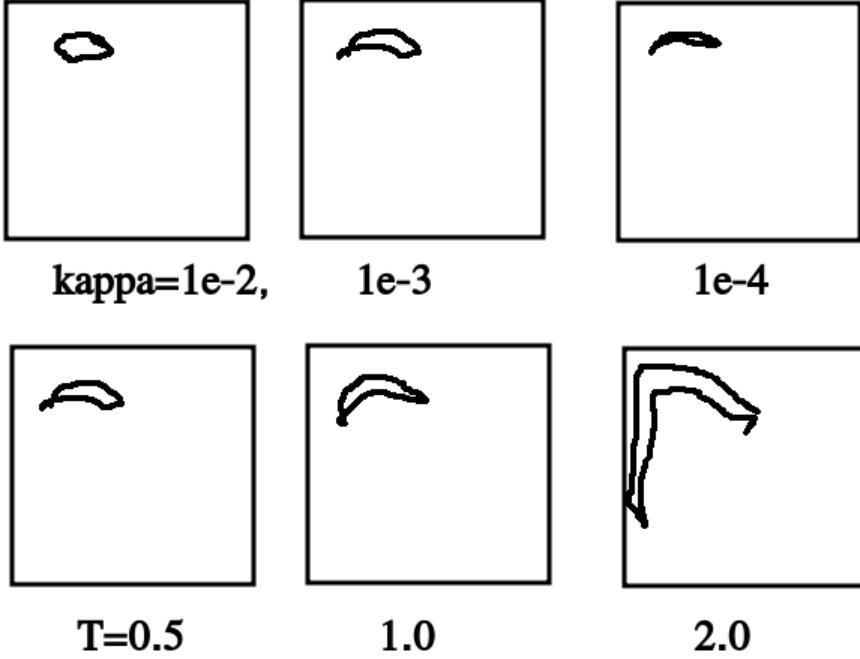


Fig. 9: Impulse responses for various diffusivities κ and fixed simulation time $T_f=0.5$ (top) and simulation times with fixed $\kappa=1e-3$ (bot). Colorbar is normalized independently for each subfigure

688

- Preconditioned spectrum, condition numbers? (1)

689

8. Conclusions. We presented an efficient matrix-free PSF method for approximating operators with locally supported non-negative integral kernels. The method only requires access to the operator via application of the operator to a small number of vectors. The idea of the method is to compute batches of impulse responses by applying the operator to Dirac combs of scattered point sources, then interpolate these impulse responses to approximate entries of the operator's integral kernel. The interpolation is based on a new principle we call "local mean displacement invariance," which generalizes and improves upon classical local translation invariance. The ability to quickly approximate arbitrary integral kernel entries permits us to form an H-matrix approximation of the operator. Fast H-matrix arithmetic is then used to perform further linear algebra operations that cannot be performed easily with the original operator, such as matrix factorization and inversion. The supports of the impulse responses are estimated to be contained in ellipsoids, which are determined a-priori via a new procedure that involves applying the operator to a small number of polynomial functions. Point source locations for the impulse response batches are chosen using a greedy ellipsoid packing procedure, in which we choose as many impulse responses per batch as possible, while ensuring that the corresponding ellipsoids do not overlap. We applied the method to approximate the Gauss-Newton Hessian in a Stokes ice sheet inverse problem, and saw that the approximation substantially outper-

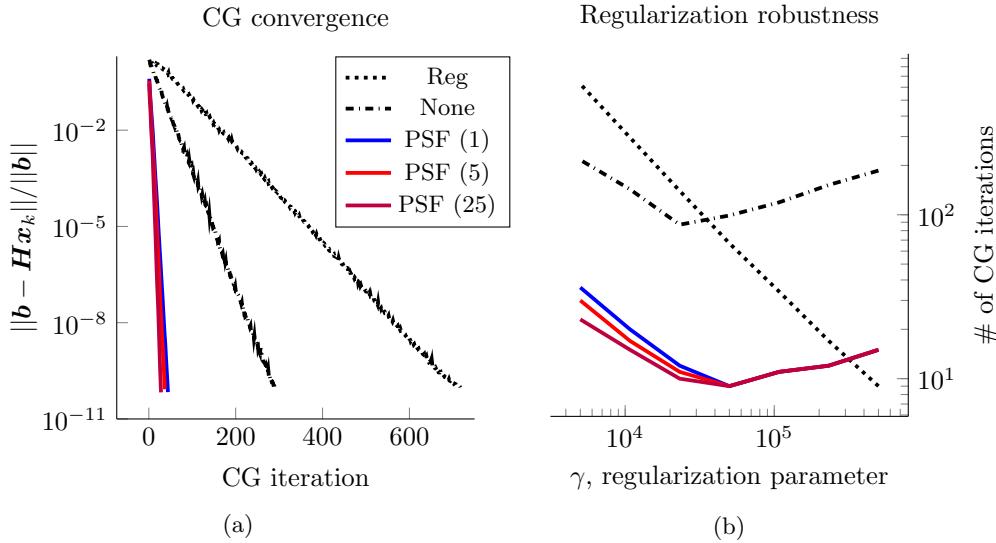


Fig. 10: Left (10a) shows the convergence history for solving $Hx = b$ using PCG, where b has i.i.d. random entries drawn from the standard normal distribution. Here we use $\gamma = 9999 \times 10^9$, which is determined by the Morozov discrepancy principle. Right (10b) shows the number of PCG iterations required to achieve $\|b - Hx_k\|/\|b\| < 10^{-6}$, where x_k is the k th PCG iterate, for a range of different κ . Results in these figures are shown for our preconditioner with 5, 15, and 30 batches (PSF (5), PSF (15), and PSF(30), respectively), regularization preconditioning (REG), and no preconditioning (NONE). The preconditioner is constructed using H_{gn} .

708 forms existing Hessian approximation methods. Although the method we presented is
 709 not applicable to all Hessians, it is applicable to many Hessians of practical interest.
 710 For these Hessians, our method offers a *data scalable* alternative to conventional low
 711 rank approximation, because our method can form high rank approximations of an
 712 operator using a small number of operator applies, and thus forward and adjoint PDE
 713 solves.

714 **Appendix A. Ellipsoid intersection test.** The procedure for choosing
 715 sample points relies on quickly determining whether two ellipsoids intersect. Let E_1
 716 and E_2 be the ellipsoids defined as

$$717 \quad E_i := \{x : (x - \mu_i)^T \Sigma_i^{-1} (x - \mu_i) \leq \tau^2\}$$

719 for $i \in \{1, 2\}$, where $\mu_1, \mu_2 \in \mathbb{R}^d$, and $\Sigma_1, \Sigma_2 \in \mathbb{R}^{d \times d}$ are positive definite. Let K be
 720 the following one dimensional convex function:

$$721 \quad K(s) := 1 - \frac{1}{\tau^2} (\mu_1 - \mu_2)^T \left(\frac{1}{1-s} \Sigma_1 + \frac{1}{s} \Sigma_2 \right)^{-1} (\mu_1 - \mu_2).$$

722 In [33] it is shown that $E_1 \cap E_2 = \{\}$ if and only if $K(s) < 0$ for some $s \in (0, 1)$.

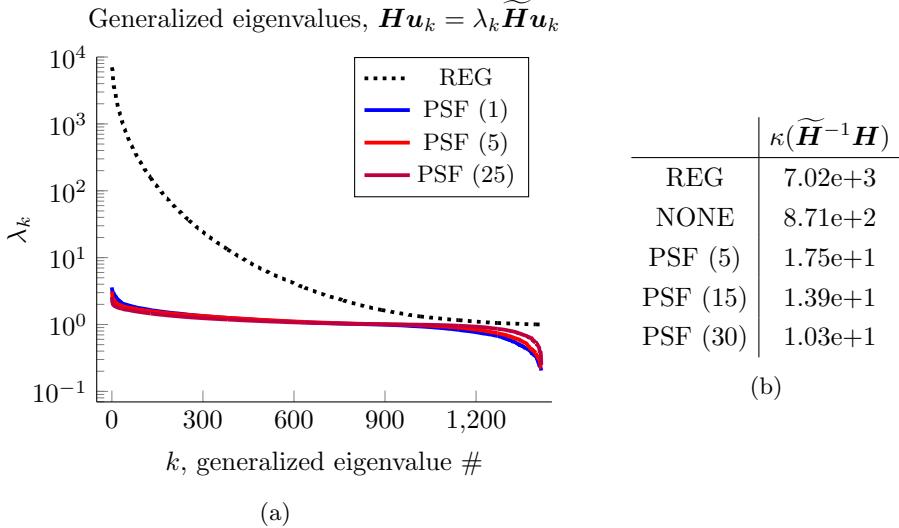


Fig. 11: Left (11a): eigenvalues for generalized eigenvalue problem $\mathbf{H}\mathbf{u}_k = \lambda_k \widetilde{\mathbf{H}}\mathbf{u}_k$. Here \mathbf{H} is the Hessian and $\widetilde{\mathbf{H}}$ is the PSF approximation constructed using the Gauss-Newton Hessian, \mathbf{H}_{gn} , for 5, 15, and 30 batches (PSF (5), PSF (15), and PSF (30), respectively), or the regularization Hessian (REG). Right (11b): condition number for $\widetilde{\mathbf{H}}^{-1}\mathbf{H}$ and $\widetilde{\mathbf{H}}^{-1}\mathbf{H}_{gn}$ for these different preconditioners, and no preconditioner (NONE). COND(H) LABELS!

723 The function $K(s)$ may be evaluated quickly for many s by computing the solution
 724 to the generalized eigenvalue problem $\Sigma_1 P = \Sigma_2 P \Lambda$, where $P \in \mathbb{R}^{d \times d}$ is the matrix
 725 of generalized eigenvectors, and $\Lambda = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_d)$ is the diagonal matrix of
 726 generalized eigenvalues λ_i . After some algebraic manipulations, we have

$$727 \quad (A.1) \quad K(s) = 1 - \frac{1}{\tau^2} \sum_{i=1}^d \frac{s(1-s)}{1 + s(\lambda_i - 1)} v_i^2,$$

728 where $v := P^T (\mu_1 - \mu_2)$. To check whether E_1 and E_2 intersect, we compute the
 729 generalized eigenvalue decomposition of Σ_1 and Σ_2 , form v , then minimize $K(s)$ in
 730 the form (A.1) on $(0, 1)$ using Brent's method [12]. If $K(s^*) < 0$ at the minimizer,
 731 s^* , then $E_1 \cap E_2 = \{\}$. Otherwise $E_1 \cap E_2 \neq \{\}$.

732 This ellipsoid intersection test is only performed if the coordinate axis aligned
 733 bounding boxes for the ellipsoids intersect. The smallest coordinate axis aligned box
 734 that contains E_i is the box $\prod_{k=1}^d [\mu_i^k - \tau \sqrt{\Sigma_i^{kk}}, \mu_i^k + \tau \sqrt{\Sigma_i^{kk}}]$, where μ_i^k is the k th
 735 entry of μ_i , and Σ_i^{kk} is the k th diagonal entry of Σ_i ⁶. The boxes $\prod_{k=1}^d [a_-^k, a_+^k]$ and
 736 $\prod_{k=1}^d [b_-^k, b_+^k]$ intersect if and only if $a_-^k \leq b_+^k$ and $b_-^k \leq a_+^k$ for $k = 1, \dots, d$.

737 **Appendix B. Rational positive semi-definite modification.** In many
 738 problems \mathcal{A} is symmetric positive semi-definite. However, \mathbf{A}_H is non-symmetric and

⁶See <https://math.stackexchange.com/q/3928964>

739 indefinite because of approximation error. This is undesirable. Symmetry and positive
 740 semi-definiteness are important properties which should be preserved if possible. Also,
 741 lacking these properties may prevent one from using highly effective algorithms, such
 742 as the conjugate gradient method, to perform further operations involving \mathbf{A}_H .

743 If \mathcal{A} is symmetric, we symmetrize \mathbf{A}_H as follows: $\mathbf{A}_H^{\text{sym}} := \frac{1}{2}(\mathbf{A}_H + \mathbf{A}_H^T)$. If
 744 \mathcal{A} is positive semi-definite, we adapt the rational method from [40, Section 13.8] to
 745 approximate the matrix absolute value of $\mathbf{A}_H^{\text{sym}}$, which we denote by $|\mathbf{A}_H^{\text{sym}}|$. That is,
 746 $|\mathbf{A}_H^{\text{sym}}|$ has the same eigenvectors as $\mathbf{A}_H^{\text{sym}}$, but the eigenvalues of $|\mathbf{A}_H^{\text{sym}}|$ are the abso-
 747 lute values of the eigenvalues of $\mathbf{A}_H^{\text{sym}}$. Once constructed, we use the approximation
 748 of $|\mathbf{A}_H^{\text{sym}}|$ in place of \mathbf{A}_H in further linear algebra operations.

749 In detail, let a be any scalar satisfying $-a < \lambda_{\min}$, where λ_{\min} is the smallest
 750 eigenvalue of $\mathbf{A}_H^{\text{sym}}$, and set

751 (B.1)
$$\mathbf{T} := (2\mathbf{A}_H^{\text{sym}} + a\mathbf{I})/a,$$

752 From the results in [40, Section 13.8] it is straightforward to show that

753 (B.2)
$$\boldsymbol{\Pi}_l := (\mathbf{I} + \mathbf{T}^{2^l})^{-1}$$

754 approximates the spectral projector onto the eigenspace of $\mathbf{A}_H^{\text{sym}}$ corresponding to
 755 negative eigenvalues. Thus, we have

756 (B.3)
$$|\mathbf{A}_H^{\text{sym}}| \approx (\mathbf{I} - 2\boldsymbol{\Pi}_l)\mathbf{A}_H^{\text{sym}}.$$

757 To form an approximation of $|\mathbf{A}_H^{\text{sym}}|$, first, we compute λ_{\min} using the Lanczos method,
 758 and set $a = \gamma\lambda_{\min}$, where $\gamma \geq 1$ is a parameter that must be chosen. The method
 759 is relatively insensitive to γ and numerically we find that the method works well for
 760 any $\gamma \in [1, 2]$. We use $\gamma = 1.5$ in our numerical results. Next, we compute \mathbf{T} per
 761 (B.1), compute \mathbf{T}^{2^l} per (B.2) via repeated squaring, compute $\boldsymbol{\Pi}_l$ per (B.2), and form
 762 an approximation of $|\mathbf{A}_H^{\text{sym}}|$ per (B.3). Fast H-matrix methods are used to perform
 763 the required matrix-matrix additions and multiplications, and matrix inversion. We
 764 recommend using $l = 1$ or 2 . Larger l is unnecessary and may make the method
 765 unstable.

766 **Acknowledgments.** We thank J.J. Alger, Longfei Gao, Mathew Hu, and Rami
 767 Nammour for helpful discussions. We thank Trevor Heise for editing suggestions. We
 768 thank Georg Stadler for domain geometry help.

- 770 [1] H.-M. ADORF, *Towards HST restoration with a space-variant PSF, cosmic rays and other*
 missing data, in The Restoration of HST Images and Spectra-II, 1994, p. 72.
- 771 [2] V. AKÇELIK, G. BIROS, A. DRĂGĂNESCU, O. GHATTAS, J. HILL, AND B. VAN BLOEMAN WAAN-
 773 DERS, *Dynamic data-driven inversion for terascale simulations: Real-time identification*
 of airborne contaminants, in Proceedings of SC2005, Seattle, 2005.
- 774 [3] A. ALEXANDERIAN, P. J. GLOOR, AND O. GHATTAS, *On Bayesian A-and D-optimal experimen-*
 776 *tal designs in infinite dimensions*, Bayesian Analysis, 11 (2016), pp. 671–695.
- 777 [4] N. ALGER, *Data-scalable Hessian preconditioning for distributed parameter PDE-constrained*
 inverse problems, PhD thesis, The University of Texas at Austin, 2019.
- 778 [5] N. ALGER, V. RAO, A. MYERS, T. BUI-THANH, AND O. GHATTAS, *Scalable matrix-free adap-*
 779 *tive product-convolution approximation for locally translation-invariant operators*, SIAM
 Journal on Scientific Computing, 41 (2019), pp. A2296–A2328.
- 780 [6] N. ALGER, U. VILLA, T. BUI-THANH, AND O. GHATTAS, *A data scalable augmented Lagrangian*
 781 *KKT preconditioner for large-scale inverse problems*, SIAM Journal on Scientific Comput-
 782 ing, 39 (2017), pp. A2365–A2393.

- [7] I. AMBARTSUMYAN, W. BOUKARAM, T. BUI-THANH, O. GHATTAS, D. KEYES, G. STADLER, G. TURKIYYAH, AND S. ZAMPINI, *Hierarchical matrix approximations of Hessians arising in inverse problems governed by PDEs*, SIAM Journal on Scientific Computing, 42 (2020), pp. A3397–A3426.
- [8] R. C. ASTER, B. BORCHERS, AND C. H. THURBER, *Parameter estimation and inverse problems*, Elsevier, 2018.
- [9] J. L. BENTLEY, *Multidimensional binary search trees used for associative searching*, Communications of the ACM, 18 (1975), pp. 509–517.
- [10] J. BIGOT, P. ESCANDE, AND P. WEISS, *Estimation of linear operators from scattered impulse responses*, Applied and Computational Harmonic Analysis, 47 (2019), pp. 730–758.
- [11] S. BÖRM, L. GRASEDYCK, AND W. HACKBUSCH, *Introduction to hierarchical matrices with applications*, Engineering analysis with boundary elements, 27 (2003), pp. 405–422.
- [12] R. P. BRENT, *An algorithm with guaranteed convergence for finding a zero of a function*, The computer journal, 14 (1971), pp. 422–425.
- [13] T. BUI-THANH, O. GHATTAS, J. MARTIN, AND G. STADLER, *A computational framework for infinite-dimensional Bayesian inverse problems Part I: The linearized case, with application to global seismic inversion*, SIAM Journal on Scientific Computing, 35 (2013), pp. A2494–A2523.
- [14] J. CHEN AND M. L. STEIN, *Linear-cost covariance functions for Gaussian random fields*, Journal of the American Statistical Association, (2021), pp. 1–18.
- [15] H. CHENG, Z. GIMBUTAS, P.-G. MARTINSSON, AND V. ROKHLIN, *On the compression of low rank matrices*, SIAM Journal on Scientific Computing, 26 (2005), pp. 1389–1404.
- [16] T. CUI, J. MARTIN, Y. M. MARZOUK, A. SOLONEN, AND A. SPANTINI, *Likelihood-informed dimension reduction for nonlinear inverse problems*, Inverse Problems, 30 (2014), p. 114015.
- [17] L. DENIS, E. THIÉBAUT, AND F. SOULEZ, *Fast model of space-variant blurring and its application to deconvolution in astronomy*, in Image Processing (ICIP), 2011 18th IEEE International Conference on, IEEE, 2011, pp. 2817–2820.
- [18] L. DENIS, E. THIÉBAUT, F. SOULEZ, J.-M. BECKER, AND R. MOURYA, *Fast approximations of shift-variant blur*, International Journal of Computer Vision, 115 (2015), pp. 253–278.
- [19] J. DUCHON, *Splines minimizing rotation-invariant semi-norms in Sobolev spaces*, in Constructive theory of functions of several variables, Springer, 1977, pp. 85–100.
- [20] H. W. ENGL, M. HANKE, AND A. NEUBAUER, *Regularization of Inverse Problems*, vol. 375 of Mathematics and Its Applications, Springer Netherlands, 1996.
- [21] C. ERICSON, *Real-time collision detection*, Crc Press, 2004.
- [22] P. ESCANDE AND P. WEISS, *Sparse wavelet representations of spatially varying blurring operators*, SIAM Journal on Imaging Sciences, 8 (2015), pp. 2976–3014.
- [23] P. ESCANDE AND P. WEISS, *Approximation of integral operators using product-convolution expansions*, Journal of Mathematical Imaging and Vision, 58 (2017), pp. 333–348.
- [24] P. ESCANDE AND P. WEISS, *Fast wavelet decomposition of linear operators through product-convolution expansions*, IMA Journal of Numerical Analysis, 42 (2022), pp. 569–596.
- [25] P. ESCANDE, P. WEISS, AND F. MALGOUYRES, *Spatially varying blur recovery. Diagonal approximations in the wavelet domain*, (2012).
- [26] A. FICHTNER AND T. v. LEEUWEN, *Resolution analysis by random probing*, Journal of Geophysical Research: Solid Earth, 120 (2015), pp. 5549–5573.
- [27] D. FISH, J. GROCHMALICKI, AND E. PIKE, *Scanning singular-value-decomposition method for restoration of images with space-variant blur*, JOSA A, 13 (1996), pp. 464–469.
- [28] H. P. FLATH, L. C. WILCOX, V. AKÇELIK, J. HILL, B. VAN BLOEMEN WAANDERS, AND O. GHATTAS, *Fast algorithms for Bayesian uncertainty quantification in large-scale linear inverse problems based on low-rank partial Hessian approximations*, SIAM Journal on Scientific Computing, 33 (2011), pp. 407–432.
- [29] P. H. FLATH, *Hessian-based response surface approximations for uncertainty quantification in large-scale statistical inverse problems, with applications to groundwater flow*, PhD thesis, The University of Texas at Austin, 2013.
- [30] M. GENTILE, F. COURBIN, AND G. MEYLAN, *Interpolating point spread function anisotropy*, Astronomy & Astrophysics, 549 (2013).
- [31] C. J. GEOGA, M. ANITESCU, AND M. L. STEIN, *Scalable gaussian process computations using hierarchical matrices*, Journal of Computational and Graphical Statistics, 29 (2020), pp. 227–237.
- [32] O. GHATTAS AND K. WILLCOX, *Learning physics-based models from data: perspectives from inverse problems and model reduction*, Acta Numerica, 30 (2021), pp. 445–554.
- [33] I. GILITSCHENSKI AND U. D. HANEBECK, *A robust computational test for overlap of two arbitrary-dimensional ellipsoids in fault-detection of Kalman filters*, in 2012 15th Inter-

- national Conference on Information Fusion, IEEE, 2012, pp. 396–401.
- [34] J. W. GLEN, *The creep of polycrystalline ice*, Proceedings of the Royal Society of London. Series A. Mathematical and Physical Sciences, 228 (1955), pp. 519–538.
- [35] L. GRASEDYCK AND W. HACKBUSCH, *Construction and arithmetics of H-matrices*, Computing, 70 (2003), pp. 295–334.
- [36] L. GRASEDYCK, R. KRIEMANN, AND S. LE BORNE, *Parallel black box \mathcal{H} -LU preconditioning for elliptic boundary value problems*, Computing and visualization in science, 11 (2008), pp. 273–291.
- [37] M. D. GUNZBURGER, *Perspectives in Flow Control and Optimization*, SIAM, Philadelphia, 2003.
- [38] W. HACKBUSCH, *A sparse matrix arithmetic based on H-matrices. Part I: Introduction to H-matrices*, Computing, 62 (1999), pp. 89–108.
- [39] W. HACKBUSCH, *Hierarchical matrices: algorithms and analysis*, vol. 49, Springer, 2015.
- [40] W. HACKBUSCH AND W. KRESS, *A projection method for the computation of inner eigenvalues using high degree rational operators*, Computing, 81 (2007), pp. 259–268.
- [41] N. HALKO, P.-G. MARTINSSON, AND J. A. TROPP, *Finding structure with randomness: Probabilistic algorithms for constructing approximate matrix decompositions*, SIAM review, 53 (2011), pp. 217–288.
- [42] T. HARTLAND, G. STADLER, M. PEREGO, K. LIEGEOIS, AND N. PETRA, *Hierarchical off-diagonal low-rank approximation of Hessians in inverse problems, with application to ice sheet model initialization*, 2023, doi:10.48550/ARXIV.2301.03644, <https://arxiv.org/abs/2301.03644>.
- [43] T. J. HUGHES, *The finite element method: linear static and dynamic finite element analysis*, Courier Corporation, 2012.
- [44] T. ISAAC, N. PETRA, G. STADLER, AND O. GHATTAS, *Scalable and efficient algorithms for the propagation of uncertainty from data through inference to prediction for large-scale problems, with application to flow of the Antarctic ice sheet*, Journal of Computational Physics, 296 (2015), pp. 348–368.
- [45] T. ISAAC, N. PETRA, G. STADLER, AND O. GHATTAS, *Scalable and efficient algorithms for the propagation of uncertainty from data through inference to prediction for large-scale problems, with application to flow of the Antarctic ice sheet*, Journal of Computational Physics, 296 (2015), pp. 348–368, doi:10.1016/j.jcp.2015.04.047.
- [46] J. KAIPIO AND E. SOMERSALO, *Statistical and Computational Inverse Problems*, vol. 160 of Applied Mathematical Sciences, Springer-Verlag New York, 2005, doi:10.1007/b138659, <https://link.springer.com/978-0-387-27132-3>.
- [47] K.-T. KIM, U. VILLA, M. PARNO, Y. MARZOUK, O. GHATTAS, AND N. PETRA, *hIPPYlib-MUQ: A bayesian inference software framework for integration of data with complex predictive models under uncertainty*, arXiv preprint arXiv:2112.00713, (2021).
- [48] R. KRIEMANN, *HLIBpro user manual*, Max-Planck-Institute for Mathematics in the Sciences, Leipzig, 48 (2008).
- [49] R. KRIEMANN, *\mathcal{H} -LU factorization on many-core systems*, Computing and Visualization in Science, 16 (2013), pp. 105–117.
- [50] P. LE TALLEC AND A. PATRA, *Non-overlapping domain decomposition methods for adaptive hp approximations of the Stokes problem with discontinuous pressure fields*, Computer Methods in Applied Mechanics and Engineering, 145 (1997), pp. 361–379.
- [51] J. LEVITT AND P.-G. MARTINSSON, *Linear-complexity black-box randomized compression of hierarchically block separable matrices*, arXiv preprint arXiv:2205.02990, (2022).
- [52] L. LIN, J. LU, AND L. YING, *Fast construction of hierarchical matrix representation from matrix–vector multiplication*, Journal of Computational Physics, 230 (2011), pp. 4071–4087.
- [53] F. LINDGREN, H. RUE, AND J. LINDSTRÖM, *An explicit link between Gaussian fields and Gaussian Markov random fields: the stochastic partial differential equation approach*, Journal of the Royal Statistical Society: Series B (Statistical Methodology), 73 (2011), pp. 423–498.
- [54] M. W. MAHONEY AND P. DRINEAS, *CUR matrix decompositions for improved data analysis*, Proceedings of the National Academy of Sciences, 106 (2009), pp. 697–702.
- [55] P.-G. MARTINSSON, *A fast randomized algorithm for computing a hierarchically semiseparable representation of a matrix*, SIAM Journal on Matrix Analysis and Applications, 32 (2011), pp. 1251–1274.
- [56] P.-G. MARTINSSON, *Compressing rank-structured matrices via randomized sampling*, SIAM Journal on Scientific Computing, 38 (2016), pp. A1959–A1986.
- [57] P.-G. MARTINSSON AND J. A. TROPP, *Randomized numerical linear algebra: Foundations and algorithms*, Acta Numerica, 29 (2020), pp. 403–572. Section 20.
- [58] J. G. NAGY AND D. P. O’LEARY, *Restoring images degraded by spatially variant blur*, SIAM

- Journal on Scientific Computing, 19 (1998), pp. 1063–1082.
- [59] J. NOCEDAL AND S. J. WRIGHT, *Numerical Optimization*, Springer, 1999.
- [60] N. PETRA, J. MARTIN, G. STADLER, AND O. GHATTAS, *A computational framework for infinite-dimensional Bayesian inverse problems, Part II: Stochastic Newton MCMC with application to ice sheet flow inverse problems*, SIAM Journal on Scientific Computing, 36 (2014), pp. A1525–A1555.
- [61] N. PETRA, H. ZHU, G. STADLER, T. HUGHES, AND O. GHATTAS, *An inexact Gauss-Newton method for inversion of basal sliding and rheology parameters in a nonlinear Stokes ice sheet model*, Journal of Glaciology, 58 (2012), pp. 889–903, doi:10.3189/2012JoG11J182.
- [62] L. ROININEN, J. M. HUTTUNEN, AND S. LASANEN, *Whittle-Matérn priors for Bayesian statistical inversion with applications in electrical impedance tomography*, Inverse Problems & Imaging, 8 (2014), p. 561.
- [63] Y. SAAD, *Iterative methods for sparse linear systems*, Society for Industrial and Applied Mathematics, Philadelphia, PA, second ed., 2003, doi:10.1137/1.9780898718003, <https://doi.org/10.1137/1.9780898718003>.
- [64] Y. SAAD AND M. SOSONKINA, *Distributed Schur complement techniques for general sparse linear systems*, SIAM Journal on Scientific Computing, 21 (1999), pp. 1337–1356.
- [65] A. SPANTINI, A. SOLONEN, T. CUI, J. MARTIN, L. TENORIO, AND Y. MARZOUK, *Optimal low-rank approximations of Bayesian linear inverse problems*, SIAM Journal on Scientific Computing, 37 (2015), pp. A2451–A2487.
- [66] A. STUART, *Inverse problems: a Bayesian perspective*, Acta Numerica, 19 (2010), pp. 451–559.
- [67] A. TARANTOLA, *Inverse Problem Theory and Methods for Model Parameter Estimation*, SIAM, Philadelphia, PA, 2005.
- [68] J. TRAMPERT, A. FICHTNER, AND J. RITSEMA, *Resolution tests revisited: the power of random numbers*, Geophysical Journal International, 192 (2013), pp. 676–680.
- [69] U. VILLA, N. PETRA, AND O. GHATTAS, *hIPPYlib: an extensible software framework for large-scale inverse problems governed by PDEs: part i: deterministic inversion and linearized Bayesian inference*, ACM Transactions on Mathematical Software (TOMS), 47 (2021), pp. 1–34.
- [70] C. R. VOGEL, *Computational Methods for Inverse Problems*, Frontiers in Applied Mathematics, Society for Industrial and Applied Mathematics (SIAM), Philadelphia, PA, 2002.
- [71] H. WENDLAND, *Scattered data approximation*, vol. 17, Cambridge University Press, 2004.
- [72] H. ZHU, S. LI, S. FOMEL, G. STADLER, AND O. GHATTAS, *A Bayesian approach to estimate uncertainty for full waveform inversion with a priori information from depth migration*, Geophysics, 81 (2016), pp. R307–R323.